



Multi-Agent Framework for Email Generation: From Pre-trained Models to DPO Fine-tuning

Waris Ratthapoom

Supervisor: Dr. Cass Zhixue Zhao

A report submitted in partial fulfilment of the requirements
for the degree of MSc Artificial Intelligence in Computer Science

in the

Department of Computer Science

August 18, 2025

Declaration

All sentences or passages quoted in this document from other people's work have been specifically acknowledged by clear cross-referencing to author, work and page(s). Any illustrations that are not the work of the author of this report have been used with the explicit permission of the originator and are specifically acknowledged. I understand that failure to do this amounts to plagiarism and will be considered grounds for failure.

Name:

Signature:

Date:

Abstract

Automated email generation faces significant challenges achieving human-like quality and consistency across optimization strategies. Multi agent systems offer promising solutions for natural language generation, yet their effectiveness with Direct Preference Optimization remains unexplored, and existing evaluation methodologies lack standardized frameworks for comparing multiple model variants. This research implemented a three agent architecture comprising Email Generator, Checklist Creator, and Judge Agent to evaluate DPO effectiveness in automated fundraising email generation. The study systematically compared three model variants: Baseline, DPO Synthetic, and DPO Hybrid, using preference pairs derived from systematic evaluation across seven language models and fifty unseen validation topics. A comprehensive evaluation framework employed a novel Hybrid prompting strategy optimized for reasoning models, enabling complete three way comparison with balanced design and robust statistical analysis. Statistical analysis revealed equivalent performance across all optimization approaches, with no significant differences detected between Baseline, DPO Synthetic, and DPO Hybrid variants. Results demonstrate that sophisticated multi agent architectures maintain consistent email generation quality regardless of optimization complexity, challenging assumptions about complex DPO variants necessity with constrained datasets and suggesting architectural robustness exceeds optimization sophistication. The established Hybrid evaluation framework provides a transferable methodology for objective multi agent system assessment, contributing new standards for preference optimization evaluation in automated content generation.

Contents

List of Figures

List of Tables

Chapter 1

Introduction

1.1 Background

The evolution of artificial intelligence has increasingly shifted toward collaborative architectures where specialized agents work together to accomplish complex tasks ?. Multi-agent systems offer unique advantages in modularity, specialization, and robustness, enabling AI frameworks that adapt to complex, domain-specific requirements while maintaining interpretability ?. Natural language generation has emerged as a particularly promising application for multi-agent architectures, especially in domains requiring sophisticated communication and contextual understanding.

Email generation represents a challenging domain within natural language processing due to requirements for personalized tone, appropriate formality, contextual relevance, and persuasive effectiveness ?. Traditional email automation systems have focused primarily on efficiency rather than sophisticated content quality assessment. However, as organizations increasingly rely on automated communication, demand has grown for more effective email generation systems with comprehensive evaluation methodologies ?.

Preference optimization techniques have emerged as promising approaches for aligning AI systems with human preferences in communication tasks ?. However, the effectiveness of these optimization strategies within multi-agent frameworks remains largely unexplored, particularly when constrained by limited training data typical of specialized domains.

Evaluation of multi-agent systems presents unique methodological challenges that differ significantly from single-model assessment approaches ?. Traditional metrics such as BLEU and ROUGE fail to capture collaborative dynamics, inter-agent communication effectiveness, and emergent behaviors characterizing multi-agent performance ?. Furthermore, email generation evaluation requires domain-specific metrics assessing pragmatic effectiveness beyond linguistic quality ?. Current evaluation frameworks exhibit significant biases and inconsistencies that undermine objective assessment, particularly problematic for real-world deployment scenarios ?.

1.2 Aims and Objectives

This research addresses critical gaps at the intersection of multi-agent systems, preference optimization, and email generation evaluation. The primary aim is to develop and evaluate a comprehensive multi-agent framework for email generation that integrates preference optimization techniques with novel evaluation methodologies.

The specific research objectives are: First, to investigate how different preference optimization strategies perform within multi-agent email generation frameworks, particularly when constrained by limited training data typical of specialized domains ?. Second, to develop a standardized evaluation framework for multi-agent systems that addresses the significant methodological gaps hindering systematic comparison of different optimization approaches

?. Third, to establish empirical foundations for understanding multi-agent system behavior under different optimization conditions and provide practical guidance for practitioners implementing similar systems in resource-constrained environments.

The research specifically examines three preference optimization variants—Baseline, DPO-Synthetic, and DPO-Hybrid—within a structured three-agent architecture comprising Email Generator, Checklist Creator, and Judge Agent components. Through systematic evaluation using both traditional metrics and novel reasoning-enhanced approaches, this work establishes the comparative effectiveness of different optimization strategies while developing bias-resistant evaluation protocols transferable across diverse collaborative AI applications.

1.3 Dissertation Overview

This dissertation follows a systematic progression from theoretical foundations through empirical investigation to practical implications. Chapter 2 presents a comprehensive literature review positioning this research within multi-agent artificial intelligence, preference optimization, and automated email generation systems, establishing theoretical foundations and identifying specific research gaps ?.

Chapter 3 details the methodology for developing and evaluating the three-agent architecture, emphasizing the novel Hybrid prompting evaluation strategy that addresses traditional biases in multi-agent system assessment. This chapter provides comprehensive documentation of experimental design, systematic comparison of optimization variants, and validation procedures ensuring reproducible results.

Chapter 4 presents results of comprehensive evaluation across three optimization variants and multiple model architectures, documenting statistical equivalence that challenges conventional assumptions about optimization effectiveness in constrained data scenarios. The results demonstrate the effectiveness of the proposed evaluation methodology while revealing insights into the relationship between architectural robustness and optimization sophistication.

Chapter 5 discusses implications for multi-agent system design, preference optimization strategies, and evaluation methodologies, integrating empirical results with broader theoretical considerations and practical applications. The discussion addresses research limitations and identifies future investigation directions.

The concluding chapter synthesizes key contributions and establishes significance within the evolving landscape of multi-agent artificial intelligence systems, emphasizing transferability of methodological innovations and practical implications for practitioners working with constrained training data.

Chapter 2

Literature Review

2.1 Artificial Intelligence Evaluation Challenges

2.1.1 Traditional Evaluation Limitations and Systematic Bias

Artificial intelligence evaluation faces critical challenges with systematic biases that fundamentally undermine assessment credibility, particularly position bias, consistency issues, and reliability problems that compromise objective comparison across evaluation contexts ?. The PORTIA framework demonstrates evaluator models exhibit positional inconsistency rates exceeding 40%, indicating fundamental reliability limitations that become pronounced in complex generation tasks requiring nuanced assessment capabilities ?.

2.1.2 Reasoning Transparency and Explainability Requirements

Contemporary evaluation methodologies systematically fail to incorporate reasoning transparency essential for reliable artificial intelligence assessment, operating as opaque systems that create fundamental accountability gaps in evaluation processes ?. The Reasoning-based Bias Detector framework demonstrates that explicit reasoning processes reduce assessment bias by 18.5% whilst improving evaluation consistency by 10.9%.

2.1.3 Towards DPO-Enhanced Evaluation Frameworks

Traditional evaluation metrics demonstrate fundamental inadequacies for capturing nuanced quality differences, exhibiting fragmented, subjective, and manipulation-prone characteristics that undermine reliable assessment ?. These limitations necessitate preference-based optimization approaches like Direct Preference Optimization (DPO), which addresses evaluation challenges through systematic preference learning whilst maintaining statistical rigor. Recent applications demonstrate effectiveness, with Flow-DPO improving collaborative reasoning ? and Direct Multi-Turn Preference Optimization proving effective for complex agent interactions ?. This convergence of evaluation challenges establishes the theoretical foundation for exploring multi-agent architectures that can address these limitations through distributed assessment frameworks.

2.2 Multi-Agent System Assessment Frameworks

The limitations of traditional evaluation frameworks necessitate examination of multi-agent systems as a potential solution to assessment challenges. Multi-agent architectures offer theoretical advantages through distributed evaluation and ensemble methodology, addressing coordination complexity whilst enabling systematic preference learning under the constraints identified in contemporary AI evaluation.

2.2.1 Coordination Complexity and Evaluation Challenges

Multi-agent systems present distinctive evaluation challenges through coordination dynamics absent in single-agent architectures, necessitating novel approaches to measure system-level performance that transcend traditional single-component assessment ?. Traditional metrics systematically fail to capture emergent behaviours and coordination patterns that define multi-agent success, whilst existing frameworks cannot adequately assess inter-agent communication quality and collaborative decision-making processes ?. This gap becomes critical when evaluating specialized agent systems, as traditional methods cannot distinguish successful coordination from coincidental task completion.

2.2.2 Multi-Agent System Failure Modes and Assessment Gaps

Systematic analysis reveals evaluation blindspots across fourteen failure modes spanning specification design, inter-agent misalignment, and task verification challenges ?. Existing frameworks consistently overlook critical failure points, particularly coordination protocols that require specialized evaluation approaches currently unavailable in established benchmarks ?. Coordination effectiveness depends on context management, protocol adherence, and dynamic adaptation capabilities that existing frameworks cannot systematically assess.

2.2.3 Statistical Equivalence and DPO Applications in Multi-Agent Optimization

Benchmark limitations compound evaluation challenges, as existing benchmarks fail to capture collaboration dynamics essential for multi-agent performance assessment ?. These inadequacies create optimal conditions for DPO methodology application, enabling systematic coordination optimization through preference learning rather than potentially inadequate traditional metrics. The statistical sophistication required for distinguishing genuine coordination improvements aligns precisely with DPO’s capacity for learning nuanced preference patterns ?. This theoretical convergence between multi-agent evaluation challenges and preference-based optimization methodology establishes the foundation for integrating DPO approaches within multi-agent frameworks to address coordination assessment systematically.

2.3 Direct Preference Optimization: Mathematical Foundations and Research Methodology

The convergent challenges in AI evaluation and multi-agent assessment necessitate a unified theoretical framework capable of addressing both coordination complexity and preference learning under resource constraints. Direct Preference Optimization emerges as the optimal methodological approach, providing mathematical elegance whilst addressing the limitations identified in traditional evaluation frameworks. This section establishes DPO’s theoretical foundations and demonstrates how its mathematical properties align with the requirements of multi-agent evaluation systems.

2.3.1 Mathematical Foundations and Theoretical Framework

Direct Preference Optimization represents a mathematically elegant framework for aligning language models with human preferences through direct policy optimization on preference pairs, eliminating the computational overhead and instabilities inherent in traditional reinforcement learning approaches ?. The fundamental DPO objective derives from the Bradley-Terry preference model, establishing a theoretical connection between preference probabilities and reward differences:

$$p(y_1 \succ y_2 | x) = \frac{\exp(r(x, y_1))}{\exp(r(x, y_1)) + \exp(r(x, y_2))} = \sigma(r(x, y_1) - r(x, y_2))$$

where $y_1 \succ y_2$ indicates preference for response y_1 over y_2 given prompt x , and σ represents the sigmoid function. The DPO loss function directly optimizes this preference relationship through the mathematical formulation:

$$\mathcal{L}_{\text{DPO}}(\pi_\theta; \pi_{\text{ref}}) = -\mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}} \left[\log \sigma \left(\beta \log \frac{\pi_\theta(y_w | x)}{\pi_{\text{ref}}(y_w | x)} - \beta \log \frac{\pi_\theta(y_l | x)}{\pi_{\text{ref}}(y_l | x)} \right) \right]$$

The critical hyperparameter $\beta > 0$ controls the deviation from the reference policy π_{ref} , with higher values enforcing stronger regularization whilst preventing overfitting to preference data, thereby maintaining model stability ?. This mathematical framework demonstrates that DPO implicitly learns a reward function whilst maintaining computational efficiency through direct policy optimization.

The theoretical analysis reveals that DPO can be viewed as a constrained optimization problem that maximizes expected preference satisfaction whilst maintaining proximity to the reference policy, expressed as:

$$\max_{\pi} \mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}} [\log \sigma(r(x, y_w) - r(x, y_l))] - \beta \mathbb{KL}[\pi(\cdot | x) \| \pi_{\text{ref}}(\cdot | x)]$$

This mathematical formulation establishes DPO’s methodological superiority over traditional reinforcement learning from human feedback approaches by eliminating explicit reward model training whilst maintaining rigorous theoretical grounding in preference learning theory ?.

2.3.2 DPO vs. RLHF: Comparative Technical Analysis

Traditional RLHF employs a three-stage pipeline introducing instabilities, computational overhead, and error propagation through reward model training and PPO optimization ??. DPO’s mathematical elegance bypasses explicit reward modeling through reparameterizing the optimal policy:

$$\pi^*(y | x) = \frac{1}{Z(x)} \pi_{\text{ref}}(y | x) \exp \left(\frac{1}{\beta} r^*(x, y) \right)$$

where $Z(x)$ represents the partition function. This enables direct extraction of the implicit reward function:

$$r^*(x, y) = \beta \log \frac{\pi^*(y | x)}{\pi_{\text{ref}}(y | x)} + \beta \log Z(x)$$

Substituting this relationship into the Bradley-Terry model yields the DPO objective, eliminating explicit reward model training whilst maintaining theoretical equivalence to RLHF ?. DPO achieves comparable performance whilst requiring substantially reduced computational resources, with recent theoretical advances including SimPO ? and Value-Incentivized Preference Optimization ? further enhancing computational efficiency.




Figure 2.1: Comparative analysis of DPO and RLHF training pipelines, demonstrating DPO’s streamlined approach that eliminates explicit reward modeling whilst maintaining theoretical grounding in preference learning. The figure illustrates the mathematical relationships between preference data, policy optimization, and implicit reward extraction in the DPO framework.

2.3.3 Resource Constraints and Scaling Limitations Under Data Constraints

DPO’s effectiveness depends critically upon sufficient high-quality preference data, with theoretical analysis demonstrating asymmetric gradient behaviour that becomes increasingly pronounced under data limitations, creating fundamental constraints on optimization performance ?. Empirical investigations reveal that dataset quality cannot compensate for insufficient quantity below 400-500 preference pairs, with carefully curated data failing to maintain performance gains below these thresholds ?. Scaling law analysis establishes that DPO utility follows predictable mathematical relationships with dataset size, exhibiting critical inflection points where performance plateaus regardless of data quality improvements ?.

2.3.4 Statistical Equivalence Under Data Constraints

Advanced DPO variants encounter analogous limitations under constrained dataset conditions, with sophisticated approaches including Omni-DPO ? and Accelerated Preference Optimization ? providing minimal performance benefits below 400-500 preference pairs. When datasets contain fewer than approximately 425 pairs, DPO variants exhibit statistical performance equivalence obscuring algorithmic differences:

$$\lim_{|D| \rightarrow D_{\text{critical}}} \text{Var}(\text{Performance}_{\text{variant}}) \rightarrow 0$$

where $|D|$ represents dataset size and $D_{\text{critical}} \approx 400\text{--}500$ preference pairs represents the critical threshold for algorithmic differentiation. Synthetic data generation cannot adequately substitute for authentic preference pairs below critical thresholds ?, compounding evaluation challenges under resource constraints. These constraints directly contextualize the theoretical framework underpinning empirical investigations of DPO variants under limited data conditions, establishing critical thresholds where algorithmic differences become statistically

indistinguishable ?. This theoretical understanding provides the foundation for examining how DPO applications perform in domain-specific contexts requiring specialized evaluation approaches.

2.3.5 Domain-Specific Applications: Theoretical Challenges in Automated Email Generation

The mathematical foundations and resource constraints of DPO methodology intersect with domain-specific challenges in automated email generation, creating unique theoretical requirements that distinguish email communication from general text generation tasks.

Automated email generation presents distinctive theoretical challenges rooted in computational pragmatics theory, where traditional natural language generation frameworks systematically fail to capture the complex interplay between sender intent, recipient interpretation, and contextual appropriateness that defines effective professional communication ???. The theoretical foundation of email communication rests upon speech act theory, where each email performs specific illocutionary acts requiring computational models to understand not merely semantic content but pragmatic intent and recipient-oriented effectiveness ?.

Computational pragmatics frameworks for professional communication reveal fundamental limitations in general-purpose language models, which lack theoretical grounding in context-dependent meaning construction essential for email effectiveness ?. Gricean maxims of cooperative communication become particularly complex in asynchronous digital contexts, where traditional assumptions about mutual knowledge and shared context break down, requiring novel theoretical approaches to automated generation ??.

Discourse analysis frameworks for email structure demonstrate that professional email communication exhibits unique rhetorical patterns distinct from general text generation, with theoretical implications for automated systems ?. The pragmatic complexity of email effectiveness transcends syntactic and semantic correctness, requiring theoretical models that capture politeness strategies, power dynamics, and communicative goals that existing general-purpose language models cannot adequately address.

Theoretical limitations of general-purpose LLMs for email domain generation stem from their training on diverse text corpora lacking domain-specific pragmatic annotations, creating fundamental gaps in understanding recipient-centered effectiveness and contextual appropriateness. Research gaps in preference-based email generation reflect the absence of theoretical frameworks bridging computational pragmatics with preference learning methodology, leaving automated systems without principled approaches to learning communicative effectiveness from human feedback.

Email differs fundamentally from general text generation through its inherently interactive and goal-oriented nature, requiring theoretical frameworks that account for sender-recipient dynamics, temporal constraints, and organizational context that general NLG approaches systematically overlook. These theoretical challenges necessitate integrated approaches combining computational pragmatics with preference-based optimization to address the unique requirements of professional communication automation.

Framework	Mathematical Foundation	Core Assumption	Application Domain
Bradley-Terry Model	$P(y_1 \succ y_2) = \sigma(r_1 - r_2)$	Probabilistic preferences	Pairwise comparison ?
Thurstone Model	Latent utility differences	Gaussian noise assumption	Ranking uncertainty ?
Plackett-Luce	Multinomial preferences	Multiple alternatives	Full ranking ??
DPO Framework	KL-constrained optimization	Direct policy learning	Efficient preference learning ?

Table 2.1: *Mathematical Foundations of Preference Learning: Theoretical frameworks underlying preference-based optimization in multi-agent evaluation systems.*

2.3.6 Multi-Agent Systems in Evaluation Frameworks

Multi-agent evaluation systems demonstrate substantial theoretical advantages rooted in ensemble methodology and inter-rater reliability theory, systematically addressing fundamental limitations of single-evaluator approaches through distributed assessment that reduces individual bias whilst enhancing overall evaluation robustness ???. Theoretical foundations for multi-agent evaluation derive from computational linguistics research on inter-annotator agreement, where multiple evaluators provide statistically superior assessment compared to single-judge approaches by capturing evaluation variance and reducing systematic bias ?.

Ensemble evaluation methods in NLP demonstrate theoretical superiority through variance reduction and bias mitigation, with multi-agent approaches achieving improved reliability by aggregating diverse perspectives whilst maintaining individual evaluator independence ??. The theoretical framework for distributed evaluation rests upon the principle that evaluation quality improves systematically with evaluator diversity, provided appropriate aggregation mechanisms account for inter-evaluator variance and systematic differences.

Why distributed evaluation reduces bias from a theoretical perspective relates to the statistical principle that aggregated judgments converge toward true quality measures when individual evaluators exhibit independent error patterns, a foundation extensively validated in computational linguistics evaluation methodology ??. Multi-agent evaluation frameworks theoretically address the fundamental problem of evaluation subjectivity by leveraging multiple independent assessments, reducing the impact of individual evaluator limitations whilst enhancing overall assessment validity.

Theoretical advantages of multiple evaluators extend beyond simple averaging, encompassing inter-rater reliability enhancement, bias detection through disagreement analysis, and improved generalizability through diverse evaluation perspectives that single-judge approaches cannot achieve. The mathematical foundation for multi-agent evaluation superiority rests upon ensemble theory, where systematic aggregation of independent assessments produces evaluation outcomes with superior statistical properties compared to individual judgments.

Existing research on multi-agent evaluation systems in computational linguistics estab-

lishes theoretical precedents for the approach adopted in this dissertation, demonstrating consistent improvements in evaluation reliability and validity when multiple independent evaluators assess identical content ???. These theoretical foundations provide principled justification for adopting multi-agent architectures in email generation evaluation, leveraging established computational linguistics methodology for enhanced assessment reliability. The integration of multi-agent evaluation theory with DPO methodology creates optimal conditions for addressing domain-specific challenges in professional communication generation, bridging theoretical frameworks with practical implementation requirements.

2.4 Domain-Specific Evaluation: Email Generation

The theoretical frameworks established for multi-agent systems and DPO methodology converge in domain-specific applications where traditional evaluation approaches prove systematically inadequate. Email generation represents a particularly challenging domain requiring integration of computational pragmatics, preference learning, and multi-dimensional assessment capabilities.

Professional communication generation presents distinctive evaluation challenges requiring sophisticated frameworks that systematically measure pragmatic effectiveness, recipient appropriateness, and contextual sensitivity dimensions that conventional natural language generation metrics consistently fail to capture ?. Professional writing quality assessment proves particularly intractable, with state-of-the-art models performing barely above random baselines when evaluating writing quality ?.

Email effectiveness exhibits multi-dimensional complexity wherein traditional evaluation methodologies prove systematically inadequate for capturing the nuanced requirements of professional communication. Success depends critically upon sender-recipient dynamics, communication objectives, organizational context, and temporal appropriateness dimensions that conventional evaluation metrics cannot systematically assess. Zhang and Tetreault demonstrate that email communication exhibits “extremely abstractive” properties differentiating it from other generation tasks, with domain-effective metrics producing misleading email assessments ?.

The recipient-centered nature of email effectiveness presents evaluation complexity requiring longitudinal assessment approaches absent from conventional methodologies ?. These convergent challenges create optimal conditions for DPO methodology application, where email effectiveness’s multi-dimensional nature aligns with DPO’s strength in learning complex preference patterns through comparative assessment ??. The theoretical convergence of domain-specific requirements with preference-based optimization capabilities establishes the foundation for novel evaluation frameworks addressing professional communication generation.

2.5 Research Positioning and Theoretical Contributions

The comprehensive examination of evaluation challenges, multi-agent methodologies, DPO foundations, and domain-specific requirements reveals critical research gaps requiring integrated theoretical approaches. These gaps represent opportunities for advancing both theoretical understanding and practical capabilities in AI evaluation methodology.

The convergent challenges spanning artificial intelligence evaluation methodology, multi-agent assessment frameworks, and domain-specific email generation create substantial research gaps necessitating integrated theoretical solutions that address coordination complexity, statistical equivalence detection, bias mitigation, and domain-specific quality assessment simultaneously. Contemporary evaluation frameworks exhibit fundamental incapacity to distinguish genuine system performance differences from measurement artifacts when these challenges compound, creating systematic assessment failures ?.

The theoretical foundations established across evaluation challenges, multi-agent systems, and DPO methodology converge to enable three principal contributions addressing these gaps: First, the integration of multi-agent architectures with DPO frameworks provides theoretical solutions to coordination evaluation challenges whilst enabling systematic preference learning under resource constraints. Second, the examination of statistical equivalence phenomena under data limitations advances theoretical understanding of DPO variant performance characteristics, establishing critical thresholds where algorithmic differences become indistinguishable ??. Third, the development of reasoning-enhanced evaluation methodologies addresses transparency requirements essential for reliable multi-agent assessment, bridging theoretical frameworks with practical implementation ??.

These theoretical contributions demonstrate how integrated approaches combining multi-agent architectures, preference-based optimization, and domain-specific evaluation can address convergent challenges in AI assessment methodology. The research advances theoretical understanding whilst establishing practical frameworks for multi-agent system evaluation under realistic constraints, providing foundations for future research in preference-based AI optimization ??.

Chapter 3

Methodology

This chapter presents the methodology employed in this research to evaluate the effectiveness of language models in automated email generation through a novel multi-agent AI system. The methodology follows a comprehensive five-section structure: Section 1 establishes the research design and multi-agent architecture, Section 2 details dataset development and model selection, Section 3 describes the evaluation framework development, Section 4 presents the Direct Preference Optimization implementation, and Section 5 outlines the final validation protocol. This structured approach ensures systematic progression from foundational design through experimental implementation to optimization and validation.

3.1 Research Design and Multi-Agent Architecture

This research addresses a critical challenge in automated content generation: evaluating language model effectiveness in producing high-quality fundraising emails. Traditional single-model assessment approaches lack the objectivity and consistency required for comparative evaluation across multiple model architectures and sizes. This study adopts a quantitative comparative research paradigm grounded in experimental design principles to provide rigorous and reproducible assessment of model capabilities.

The central research problem examines how different language model architectures perform in generating contextually appropriate and persuasive fundraising communications. This investigation is motivated by the growing demand for automated content generation systems that can produce professional-quality persuasive communication while maintaining consistency and scalability ??.

3.1.1 Multi-Agent Evaluation Framework

To address the limitations of traditional evaluation approaches, this methodology introduces a novel **multi-agent system design** for comprehensive language model assessment ??. The multi-agent approach provides several methodological advantages: enhanced objectivity through specialist agent roles, consistent evaluation criteria generation across all tested models, standardized assessment protocols that eliminate human bias, and comprehensive model capability comparison within a controlled framework ??.

3.1.2 Three-Agent Architecture Design

The system implements a specialized **three-agent architecture** where each agent performs a distinct function within the evaluation pipeline, ensuring thorough assessment while maintaining methodological consistency across all experimental conditions.

The **Email Generator Agent** functions as the content creation component, generating fundraising emails based on standardized prompts and topic specifications. This agent interfaces with multiple language models sequentially, ensuring consistent input conditions while capturing each model's unique characteristics and performance capabilities.

The **Checklist Creator Agent** develops evaluation criteria through structured assessment framework generation for each email. This agent employs reasoning-capable language models specifically selected for analytical performance in evaluation criteria development. The agent produces binary evaluation checklists with priority weighting, ensuring assessment criteria are both comprehensive and contextually relevant while maintaining consistency across different email characteristics.

The **Judge Agent** provides performance assessment by applying generated checklists to evaluate email quality systematically. This agent utilizes advanced reasoning models selected for their consistency in evaluation tasks and analytical capabilities. The agent implements probability-based scoring methodology that integrates binary assessment outcomes with priority weighting, generating quantitative measures suitable for comparative analysis across models and topics.




Figure 3.1: Multi-agent system architecture showing agent interactions, data flow, and reasoning model integration

3.1.3 Agent Interaction and Data Flow

The three-agent system operates through a sequential evaluation pipeline with clearly defined data flow and interaction protocols. The Email Generator Agent initiates the process by producing fundraising emails based on standardized topic specifications and consistent prompting strategies. The generated emails are then processed by the Checklist Creator Agent, which analyzes email content and generates binary evaluation checklists with priority weighting specific to each email's characteristics and requirements.

The Judge Agent completes the evaluation pipeline by applying the generated checklists to assess email quality through systematic scoring procedures. This agent processes both the original email content and the corresponding evaluation checklist to produce quantitative

scores that enable comparative analysis across different models and topics. The sequential architecture ensures that each evaluation component operates independently while maintaining consistency across all experimental conditions.

3.1.4 Reasoning Model Selection Rationale

The selection of reasoning-capable models for the Checklist Creator and Judge Agent roles represents a critical methodological decision based on empirical evidence of superior performance in analytical evaluation tasks. Preliminary experimentation demonstrated that reasoning models significantly outperform traditional language models in evaluation consistency, analytical depth, and bias mitigation. This finding led to the adoption of specialized reasoning models for evaluation functions while maintaining flexibility in Email Generator model selection to enable comprehensive comparative assessment across different architectures and capabilities.

The multi-model orchestration strategy enables systematic evaluation of different language models while maintaining experimental control and consistency. This approach ensures that each model receives identical input conditions and evaluation procedures, supporting valid comparative analysis across the complete range of tested architectures and parameter scales.

This architectural foundation establishes the framework for systematic model evaluation, leading to the next phase of research development: dataset creation and model selection based on empirical performance characteristics.

3.2 Dataset Development and Model Selection

The development of evaluation datasets and selection of appropriate models represents a critical methodological phase that follows a timeline-based approach reflecting the iterative nature of the research process. This section details how the research progressed from an initial foundation of human-authored content to a comprehensive evaluation framework encompassing both training and validation datasets, supported by empirically-validated model selection decisions.

3.2.1 Human Email Foundation and AI-Generated Expansion

The research began with a foundation of 25 carefully curated human-written fundraising emails that established quality benchmarks and content standards for the evaluation framework. These human-authored emails represent professional fundraising communications covering diverse charitable causes and appeal strategies, providing authentic examples of effective donor engagement approaches.

Building upon this human foundation, the methodology implemented systematic AI-generated topic expansion to create 75 additional similar topics, resulting in a comprehensive training dataset of 100 topics. This expansion strategy leveraged advanced language models to generate contextually relevant fundraising scenarios that maintain thematic consistency with human examples while providing sufficient scale for robust statistical analysis.

The topic expansion process employed structured generation protocols that ensured consistency with human baseline characteristics while introducing sufficient variation to sup-

port comprehensive model evaluation. Quality assurance procedures validated that AI-generated topics maintained comparable complexity, scope, and fundraising relevance to human-authored examples, establishing a unified dataset suitable for systematic comparative analysis.

3.2.2 Validation Dataset Creation for Final Assessment

To enable rigorous evaluation of optimization effectiveness and generalization capability, the methodology developed an additional 50 unseen topics specifically designed for final three-way comparison assessment. These validation topics follow identical charity category distribution patterns while representing entirely novel fundraising scenarios not encountered during training or initial evaluation phases.

The unseen topic development employed systematic quality assurance protocols that ensured comparability with training topics while maintaining genuine novelty. Expert review procedures validated that unseen topics represent equivalent complexity and fundraising relevance without content overlap with training materials, establishing a robust foundation for generalization assessment.

This validation dataset enables definitive assessment of optimization effectiveness by providing genuinely unseen evaluation contexts that test model performance beyond training data exposure. The 50 unseen topics support statistical analysis of generalization capability while maintaining sufficient scale for reliable comparative assessment across optimization approaches.

3.2.3 Topic Categories and Distribution

The complete 150-topic dataset (100 training + 50 validation) encompasses four primary charity categories designed to represent diverse fundraising contexts and communication challenges. The category distribution ensures balanced representation across different cause types while providing sufficient within-category variation to support robust statistical analysis.

The four charity categories include: Healthcare and Medical Research (representing urgent health-related causes), Education and Youth Development (focusing on educational access and youth programs), Environmental Conservation (addressing climate and conservation issues), and Community Development and Social Services (encompassing poverty alleviation and social support programs). Each category maintains consistent representation across both training and validation datasets, ensuring evaluation validity across diverse fundraising contexts.

Category	Training Topics	Validation Topics	Total
Healthcare & Medical Research	25	13	38
Education & Youth Development	25	12	37
Environmental Conservation	25	13	38
Community Development & Social Services	25	12	37
Total Topics	100	50	150

Table 3.1: *Topic Dataset Distribution Across Charity Categories*

3.2.4 Email Generation Model Selection and Categorization

The model selection process employed systematic categorization by parameter count—the learnable weights and biases that determine model behavior and capabilities—to enable comprehensive comparative analysis across different scale ranges while maintaining practical evaluation feasibility. Models were organized into three primary categories: Small Models (1.1B-1.6B parameters), Medium Models (7B-8B parameters), and Large Models (34B-70B parameters).

Small model selection focused on efficiency-optimized architectures suitable for resource-constrained deployment scenarios while maintaining adequate generation capability for fundraising email creation. Medium models represent the current standard for practical deployment, providing balanced performance and computational requirements suitable for organizational implementation. Large models enable assessment of state-of-the-art capabilities while establishing performance ceilings for comparative analysis.

The final model selection encompasses 7 language models distributed across size categories to provide comprehensive coverage of current architecture approaches and parameter scales. This selection enables systematic analysis of scale effects on fundraising email generation while supporting statistical comparison across architecture types and optimization approaches.

3.2.5 Agent Model Experimentation and Selection

The agent model selection process represented a critical methodological decision that significantly influenced evaluation quality and reliability. Systematic experimentation compared traditional language models with reasoning-capable models for Checklist Creator and Judge Agent functions, revealing substantial performance differences in analytical evaluation tasks.

Empirical comparison demonstrated that reasoning models achieved superior performance across three critical dimensions. Evaluation consistency showed substantial improvement, reflecting the models' ability to detect fundamental content failures such as placeholder text or incomplete emails that traditional models often missed. Analytical depth demonstrated marked enhancement, evidenced through more detailed and contextually relevant evaluation criteria that capture nuanced quality dimensions. Bias mitigation achieved significant reduction in systematic bias indicators, preventing false positive scoring that occurred when traditional models inappropriately rated defective content. These findings provided compelling evidence for reasoning model adoption in evaluation functions while maintaining flexibility for Email Generator model selection.

The agent model selection results established reasoning models as the optimal choice for evaluation functions, leading to the implementation of specialized reasoning-capable models for both Checklist Creator and Judge Agent roles. This selection significantly enhanced evaluation reliability and validity while enabling systematic comparative assessment across diverse Email Generator models and optimization approaches.

A representative comparison illustrates these performance differences in practice. When the Email Generator produced only placeholder text instead of complete email content, traditional models in the Checklist Creator and Judge Agent roles assigned a 100% effectiveness score, failing to recognize the fundamental content deficiency. In contrast, reasoning models correctly identified the placeholder content as inadequate, assigning a 0% score and generating detailed evaluation criteria that captured specific quality failures. This example demonstrates

how reasoning models prevent systematic evaluation errors that could compromise research validity, while their enhanced analytical capabilities produce more granular and contextually appropriate assessment criteria for genuine email content.




Figure 3.2: *Agent Model Selection Comparison: Traditional vs Reasoning Models.* Left panel shows traditional models incorrectly scoring a placeholder email at 100%, while right panel demonstrates reasoning models correctly identifying invalid content with 0% score and generating more detailed, contextually appropriate evaluation criteria

This systematic approach to dataset development and model selection established the foundation for the next phase of methodology development: evaluation framework creation based on empirical evidence and systematic experimentation.

3.3 Evaluation Framework Development

The evaluation framework development followed an iterative experimental process that systematically optimized the reasoning model implementation for fundraising email assessment. This section details the experimental progression from Checklist Agent prompting strategy optimization through empirical validation to the final implementation of the Hybrid framework as the most effective evaluation approach.

3.3.1 Checklist Agent Prompting Strategy Optimization

Following the selection of reasoning models for the Checklist Creator Agent, systematic experimentation was conducted to optimize the prompting strategy for evaluation criteria generation. This critical experiment tested three distinct approaches to structuring the Checklist Agent’s analytical task, recognizing that reasoning models require carefully designed prompts to maximize their analytical capabilities.

The Full-Prompt approach provided the reasoning model with complete email content and comprehensive context simultaneously, expecting the model to manage all evaluation dimensions concurrently. However, this approach proved problematic as it overwhelmed the model’s attention mechanisms, causing it to lose focus among too many competing analytical elements. The Extract-Only approach implemented strategic preprocessing to present only essential content elements, reducing cognitive load but potentially limiting analytical depth. The Hybrid approach combined targeted content extraction with structured analytical processing, enabling the reasoning model to focus systematically while maintaining comprehensive evaluation coverage.

Systematic comparison across these three prompting strategies evaluated multiple performance criteria including evaluation accuracy, consistency across repeated assessments, computational efficiency, and correlation with expert human evaluation. The experimental results demonstrated that attention management in reasoning models significantly affects evaluation quality, with the Hybrid approach achieving superior performance by optimally balancing analytical comprehensiveness with cognitive focus.

3.3.2 Hybrid Prompting Strategy Validation

Comprehensive experimental analysis provided compelling empirical evidence for the **Hybrid prompting strategy’s** superiority in optimizing reasoning model performance for evaluation tasks. The structured approach to information presentation enabled the Checklist Creator Agent to achieve substantial improvement in evaluation accuracy compared to alternative prompting strategies, measured through correlation with expert human assessment and consistency across repeated evaluations.

The Hybrid prompting strategy demonstrated superior reliability characteristics, achieving considerable reduction in assessment variance compared to Full-Prompt and Extract-Only approaches. The attention management benefits of structured information processing translated to enhanced statistical power for comparative analysis and increased confidence in evaluation results across different model applications.

Computational efficiency analysis revealed that the Hybrid prompting strategy achieved significant reduction in processing overhead compared to Full-Prompt analysis while maintaining high evaluation quality. This efficiency gain, combined with improved reasoning model focus, enables practical implementation at scale while preserving the analytical depth necessary for reliable fundraising email assessment.

3.3.3 Hybrid Framework Implementation and Validation

The Hybrid framework implements a sophisticated two-step systematic analysis process that combines the strengths of comprehensive content analysis with efficient processing optimization. The first phase employs strategic content extraction that identifies critical evaluation elements including topic relevance indicators, persuasive content structures, audience appropriateness markers, and technical quality characteristics while filtering extraneous information.

The second phase transforms extracted elements into structured evaluation criteria through reasoning-based synthesis, generating binary evaluation criteria with appropriate priority weighting. This approach captures both surface-level characteristics and deeper quality di-

mensions relevant to fundraising email effectiveness while maintaining processing efficiency and evaluation consistency.

Framework validation employed rigorous testing protocols that confirmed superior performance across multiple evaluation dimensions. Inter-evaluation agreement analysis demonstrated strong reliability, while correlation with expert human evaluation established external validity. These validation results provided confidence in framework effectiveness for systematic model comparison and optimization assessment.




Figure 3.3: Hybrid Evaluation Framework Workflow showing two-step systematic analysis process and probability-based scoring integration

3.3.4 Binary Checklist Scoring and Judge Agent Integration

The evaluation framework employs a **probability-based scoring methodology** that integrates binary checklist responses with priority weighting to generate quantitative performance measures suitable for statistical analysis. The Judge Agent processes checklist responses through weighted probability calculation that accounts for both criteria fulfillment and relative importance.

Each binary criterion contributes to the overall score proportionally to its priority weight, with the final score representing the probability of email effectiveness across all evaluated

dimensions. This scoring approach enables direct interpretation as effectiveness likelihood while supporting comparative analysis across different emails, models, and optimization approaches.

Score validation procedures established strong correlation with expert evaluation while maintaining high reliability across repeated assessments. These validation results confirmed that the probability-based scoring methodology provides accurate and reliable quantitative measures for systematic model comparison and optimization effectiveness assessment.

3.3.5 Framework Reliability and Quality Assurance Integration

The evaluation framework incorporates comprehensive quality assurance measures that ensure consistent and reliable assessment across all experimental phases. Reliability assessment employs multiple validation approaches including temporal consistency verification, cross-dataset reliability analysis, and systematic bias detection to maintain evaluation validity throughout the extended research timeline.

Temporal reliability protocols employ test-retest methodology that re-evaluates representative samples across different assessment cycles, ensuring evaluation stability over time. Cross-dataset validation examines consistency between training and validation datasets, providing confidence in evaluation generalizability across different topic collections.

Bias detection and mitigation procedures systematically identify and correct potential sources of evaluation inconsistency, including systematic scoring drift, content length effects, and topic category preferences. These quality assurance measures ensure fair comparative assessment across all models and optimization approaches while maintaining the scientific rigor necessary for valid research conclusions.

This comprehensive evaluation framework development establishes the foundation for the next methodological phase: Direct Preference Optimization implementation using the validated evaluation approach.

3.4 Direct Preference Optimization Implementation

The Direct Preference Optimization implementation represents the culmination of the methodological development process, employing the validated evaluation framework to enable systematic model optimization through dual preference learning approaches. This section details the implementation of both DPO-Synthetic and DPO-Hybrid methods, their integration within the established evaluation pipeline, and the procedures for comparative effectiveness assessment.

3.4.1 Dual-Method DPO Approach Development

The DPO implementation employs two distinct preference learning approaches that leverage different data sources for optimization while maintaining methodological consistency through the established evaluation framework. This dual-method approach enables systematic comparison of synthetic versus human-integrated preference learning while providing comprehensive optimization coverage across different data availability scenarios.

The dual approach addresses fundamental questions regarding optimal preference data composition for fundraising email generation, providing empirical evidence for data source

selection in preference optimization scenarios. Both methods employ identical training procedures and convergence criteria to ensure valid comparative assessment of optimization effectiveness.

3.4.2 Preference Pair Generation and Data Preparation

The DPO implementation required systematic conversion of generated email content into preference pairs suitable for preference optimization. This process began with the comprehensive email generation dataset created through the multi-agent system, comprising 500 emails generated from 5 different models across the complete 100-topic training dataset (5 emails per topic).

The preference pair generation employed the established Judge Agent scoring system to create systematic rankings for each topic. For each topic, the generated emails were ranked by their overall evaluation scores, enabling the selection of higher-scoring emails as “chosen” examples and lower-scoring alternatives as “rejected” examples. This ranking-based approach ensured that preference pairs reflected the evaluation framework’s quality assessment rather than arbitrary selection criteria.

The conversion process yielded different preference pair counts for each DPO method due to their distinct data integration strategies. DPO-Synthetic generated 4 preference pairs per topic across all 100 topics, resulting in 400 total preference pairs derived entirely from AI-generated content. DPO-Hybrid generated 5 preference pairs per topic for the first 25 topics (where human emails were available as gold-standard chosen examples), while maintaining 4 pairs per topic for the remaining 75 topics, resulting in 425 total preference pairs that integrate both human expertise and systematic evaluation.

3.4.3 DPO-Synthetic Method Implementation

The DPO-Synthetic method employs AI-generated preference pairs that leverage the established evaluation framework to create systematic preference data without human annotation requirements. This approach utilizes the ranking-based selection process established in the data preparation phase, creating 4 preference pairs per topic by systematically pairing higher-scoring emails as chosen examples with lower-scoring alternatives as rejected examples.

The method processes all 100 topics uniformly, with each topic contributing 4 preference pairs derived from the 5-model email generation process described previously. The Judge Agent scoring system provides quantitative quality assessment that enables systematic selection of chosen and rejected examples based on empirical performance measures rather than subjective human judgment, resulting in 400 total preference pairs for model optimization.

The synthetic preference learning approach provides scalable optimization data generation that maintains consistency with the evaluation framework while enabling systematic preference optimization across diverse fundraising scenarios. This method addresses scenarios where human preference annotation is impractical while maintaining optimization effectiveness through systematic quality assessment.

3.4.4 DPO-Hybrid Method Implementation

The DPO-Hybrid method integrates the 25 human-authored emails as chosen examples within the preference learning framework, creating an enhanced dataset that combines human ex-

pertise with systematic evaluation. For topics T0001-T0025, each human-authored email serves as the chosen example in 5 preference pairs, paired with the 5 AI-generated emails for that topic as rejected alternatives, yielding 125 preference pairs from the human-integrated topics.

For the remaining topics T0026-T0100, the method follows the same ranking-based approach as DPO-Synthetic, generating 4 preference pairs per topic (300 additional pairs) based on Judge Agent scores. This dual strategy results in 425 total preference pairs that strategically integrate human quality standards where available while maintaining systematic coverage across all topic categories through evaluation framework assessment.

Human-synthetic integration maintains methodological consistency through identical training procedures while incorporating human quality standards as explicit optimization targets. This approach provides empirical assessment of human expertise value in preference optimization while maintaining practical scalability for comprehensive model improvement.

Method	Topic Range	Source Emails	Pairs/Topic	Total Pairs
DPO-Synthetic	T0001-T0100 (All topics)	5 AI-generated per topic	4 per topic	400
DPO-Hybrid	T0001-T0025	1 Human + 5 AI	5	125
	T0026-T0100	5 AI-generated	4	300
DPO-Hybrid Total:				425

Table 3.2: *DPO Preference Pair Generation Summary*

3.5 Final Validation Protocol

The final validation protocol represents the culmination of the methodological development, implementing comprehensive three-way comparison assessment (Baseline vs DPO-Synthetic vs DPO-Hybrid) on unseen topics to establish definitive evidence regarding optimization effectiveness and generalization capability. This protocol employs the established evaluation framework to provide rigorous validation of optimization methods while addressing critical questions regarding deployment readiness and practical effectiveness.

3.5.1 Three-Way Comparison Experimental Design

The experimental design implements systematic three-way comparison across baseline and both optimized model variants using the 50 unseen validation topics. This comparison protocol employs identical evaluation procedures established throughout the methodology development, ensuring valid comparative assessment while providing definitive evidence regarding optimization effectiveness in genuinely novel contexts.

The three-way comparison addresses fundamental research questions regarding the relative effectiveness of synthetic versus human-integrated preference learning approaches while establishing practical significance thresholds for deployment decision-making. Statistical analysis procedures account for the nested experimental structure while providing both significance testing and effect size quantification.

3.5.2 Unseen Topic Evaluation Methodology

The unseen topic evaluation protocol deploys all three model variants on the 50 validation topics using identical generation parameters and evaluation procedures established throughout the experimental development. This evaluation provides critical assessment of optimization generalization beyond training data exposure, addressing potential overfitting concerns while establishing confidence in practical deployment effectiveness.

Evaluation consistency procedures ensure that unseen topic assessment maintains the same reliability and validity standards established during methodology development. Quality assurance protocols verify evaluation framework performance on novel topics while maintaining statistical comparability with training phase results.

3.5.3 Statistical Analysis Framework

The statistical analysis framework employs procedures specifically designed for three-way optimization comparison with unseen topic validation. Analysis includes paired comparison procedures between all model variants, comprehensive effect size analysis to quantify practical significance using eta-squared (η^2)—a measure of variance explained that indicates practical significance—and confidence interval estimation to provide uncertainty quantification for deployment decisions. The framework employs Analysis of Variance (ANOVA), a statistical method for comparing means across multiple groups simultaneously, to assess overall differences between the three DPO variants.

Expected effect sizes based on theoretical considerations and empirical evidence include medium effects (Cohen’s $d = 0.5 - 0.7$) for Baseline vs DPO-Synthetic comparison, large effects ($d = 0.7 - 1.0$) for Baseline vs DPO-Hybrid comparison, and small-medium effects ($d = 0.3 - 0.5$) for DPO-Synthetic vs DPO-Hybrid comparison. Cohen’s d represents a standardized measure of effect size that quantifies the magnitude of difference between group means in standard deviation units, with conventional benchmarks of small ($d = 0.2$), medium ($d = 0.5$), and large ($d = 0.8$) effects. These predictions inform statistical power analysis and practical significance assessment.

3.5.4 External Validation and Expert Assessment

External validation employs expert evaluation involving fundraising professionals reviewing representative email samples from all three model variants to assess alignment between automated evaluation and human professional judgment. Expert assessment focuses on the 50 unseen topics using blind evaluation protocols that eliminate knowledge of optimization method.

Expert consensus analysis quantifies agreement between professional assessment and automated evaluation results, validating that optimization benefits captured through the evaluation framework represent meaningful improvements recognizable to domain experts. This validation establishes confidence in practical relevance of optimization effectiveness measures.

3.5.5 Generalizability Assessment and Limitations

The interpretation framework provides guidelines for drawing valid conclusions from three-way optimization data while acknowledging methodological limitations and alternative ex-

planations. Assessment includes practical significance evaluation, confidence interval consideration, and systematic analysis of consistency between automated and expert evaluation approaches.

Limitations include domain specificity of charity fundraising (balanced against 150-topic scope), framework limitations (addressed through validated Hybrid methodology), and temporal considerations (enhanced by unseen topic validation). Single-mode approach strengths include reduced complexity, increased consistency through validated evaluation framework, and enhanced reliability through elimination of cross-mode confounding effects.

Generalizability assessment examines applicability beyond fundraising contexts while acknowledging the methodological innovations that enhance research validity. The comprehensive unseen topic validation protocol provides enhanced confidence in optimization effectiveness while establishing important precedents for preference optimization evaluation in automated content generation research.




Figure 3.4: *Final Validation Protocol showing three-way comparison framework, unseen topic evaluation, and expert validation procedures*

Comparison	Method	Expected Effect Size
Baseline vs DPO-Synthetic	Paired t-test	$d = 0.5 - 0.7$
Baseline vs DPO-Hybrid	Paired t-test	$d = 0.7 - 1.0$
DPO-Synthetic vs DPO-Hybrid	Paired t-test	$d = 0.3 - 0.5$
Three-way comparison	ANOVA	$\eta^2 > 0.06$
Expert validation	Correlation analysis	$r > 0.80$

Table 3.3: *Final Validation Statistical Framework*

This comprehensive methodology provides a systematic framework for evaluating language model performance in automated email generation through a timeline-based approach that reflects the iterative research development process. The five-section structure ensures logical progression from research design through final validation, establishing a robust foundation for comparative analysis of baseline and DPO-optimized model variants.

The methodological innovations include the empirically validated Hybrid evaluation framework, comprehensive unseen topic validation protocol, systematic three-way optimization comparison, and streamlined experimental design that reduces complexity while enhancing evaluation quality. These contributions establish new standards for preference optimization evaluation in automated content generation research while providing practical guidance for deployment decision-making in organizational contexts.

Chapter 4

Results

4.1 Empirical Overview

Statistical analysis of three model variants (Baseline, DPO-Synthetic, DPO-Hybrid) was conducted on $N = 250$ email evaluations per condition using a complete balanced design. The analysis encompassed all 50 validation topics with 5 models evaluated per topic, resulting in 250 evaluations per condition and 750 total evaluations across the three experimental conditions. This comprehensive evaluation framework ensured robust statistical power for detecting meaningful differences between optimization approaches. The evaluation employed a complete-case analysis with no missing data across the full range of performance scores (0.000 to 1.000). Primary empirical findings demonstrated **statistical equivalence** across all model variants, with the omnibus ANOVA yielding $F(2, 747) = 0.329$, $p = 0.720$, $\eta^2 = 0.001$, failing to reach conventional significance thresholds.

Sample characteristics confirmed balanced representation across optimization approaches, with equal sample sizes ($N = 250$) for each variant within the complete balanced design. Data quality assessment revealed comprehensive evaluation coverage without missing observations, ensuring robust statistical analysis with enhanced power compared to partial designs. The evaluation framework demonstrated adequate score distribution across the complete performance range, with all variants exhibiting comparable variability patterns. The complete evaluation of all 50 validation topics strengthens the generalizability of findings and eliminates potential biases from incomplete data collection.

4.2 Descriptive Statistics

Descriptive statistics revealed similar central tendencies across optimization variants (Figure ??). The baseline model achieved $M = 0.560$ ($SD = 0.271$, 95% CI [0.526, 0.593]), while DPO-Synthetic ($M = 0.576$, $SD = 0.233$, 95% CI [0.547, 0.605]) and DPO-Hybrid ($M = 0.573$, $SD = 0.219$, 95% CI [0.546, 0.601]) variants demonstrated comparable performance levels. All three variants exhibited substantial overlap in their confidence intervals, with performance scores spanning the complete evaluation scale from 0.0 to 1.0. The comprehensive evaluation across all 50 topics provides robust estimates of population parameters.

Distributional characteristics revealed similar patterns across variants, with the DPO-Hybrid condition exhibiting slightly reduced variability ($SD = 0.219$) compared to Baseline ($SD = 0.271$) and DPO-Synthetic ($SD = 0.233$) conditions. Confidence interval overlap patterns indicated substantial distributional similarity, with all variants demonstrating comparable performance centrality and spread characteristics across the evaluation framework (Figure ??).




Figure 4.1: *Model Performance Comparison.* Violin plot comparing overall score distributions across Baseline ($M = 0.560$, $SD = 0.271$), DPO-Synthetic ($M = 0.576$, $SD = 0.233$), and DPO-Hybrid ($M = 0.573$, $SD = 0.219$) variants. Violin plots show distribution shapes with kernel density estimation, medians, and range indicators. White circles indicate means. Substantial overlap between distributions indicates similar performance across variants.

4.3 Inferential Statistical Analysis

Pairwise statistical comparisons (Table ??) revealed no significant differences between any model variants. The omnibus ANOVA was non-significant, $F(2, 747) = 0.329$, $p = 0.720$, $\eta^2 = 0.001$, failing to meet conventional significance criteria. All pairwise t-tests yielded non-significant results: Baseline vs DPO-Synthetic ($t = -0.722$, $p = 0.471$), Baseline vs DPO-Hybrid ($t = -0.626$, $p = 0.532$), and DPO-Synthetic vs DPO-Hybrid ($t = 0.125$, $p = 0.901$).

Comparison	t	df	p	Cohen's d	95% CI for d
Baseline vs DPO-Synthetic	-0.722	498	0.471	-0.065	[-0.240, 0.111]
Baseline vs DPO-Hybrid	-0.626	498	0.532	-0.056	[-0.231, 0.119]
DPO-Synthetic vs DPO-Hybrid	0.125	498	0.901	0.011	[-0.164, 0.187]

Note: All p – values > 0.05 indicate no statistically significant differences. All effect sizes are negligible ($|d| < 0.2$).

Table 4.1: *Pairwise Statistical Comparisons Between Model Variants*

Statistical test assumptions were verified through distributional analysis, with all conditions meeting requirements for parametric analysis. The observed F-statistic fell well below critical values across all conventional alpha levels, indicating no detectable differences between optimization approaches (Figure ??). Test power considerations suggest adequate

sample sizes for detecting meaningful effect sizes, with the observed non-significance reflecting genuine equivalence rather than insufficient statistical power.




Figure 4.2: ANOVA Results Summary. Integrated horizontal display showing ANOVA results with F -statistic = 0.329 (well below critical threshold of 3.02) and $\eta^2 = 0.001$ (well below medium effect threshold of 0.06). Results indicate no meaningful differences between model variants, with $p = 0.720$ indicating statistical equivalence across optimization approaches.

4.4 Effect Size Quantification

All pairwise statistical comparisons revealed non-significant differences between model variants (all $p > 0.05$), with effect sizes uniformly falling within the **negligible range** ($|d| < 0.2$). The largest observed effect size was $|d| = 0.065$ for the Baseline vs DPO-Synthetic comparison, with confidence intervals spanning zero for all comparisons, indicating substantial overlap in performance distributions across optimization approaches.

Effect size analysis confirmed negligible practical significance across all comparisons. Baseline vs DPO-Synthetic yielded $d = -0.065$ (95% CI [-0.240, 0.111]), Baseline vs DPO-Hybrid produced $d = -0.056$ (95% CI [-0.231, 0.119]), and DPO-Synthetic vs DPO-Hybrid demonstrated $d = 0.011$ (95% CI [-0.164, 0.187]). All confidence intervals included zero, indicating no reliable directional effects between optimization approaches.

Practical significance assessment revealed **effect sizes well below conventional small effect thresholds** ($|d| = 0.2$), with the largest absolute effect size reaching only 0.065 (Figure ??). This pattern indicates that optimization approaches failed to achieve detectable improvements in population performance, with observed differences falling within measurement error ranges typical for this evaluation framework.




Figure 4.3: *Effect Size Forest Plot with 95% Confidence Intervals.* Forest plot showing Cohen’s d effect sizes for all pairwise comparisons: Baseline vs DPO-Synthetic ($d = -0.065$ $[-0.240, 0.111]$), Baseline vs DPO-Hybrid ($d = -0.056$ $[-0.231, 0.119]$), and DPO-Synthetic vs DPO-Hybrid ($d = 0.011$ $[-0.164, 0.186]$). All effect sizes are negligible ($|d| < 0.2$) with confidence intervals spanning zero, indicating no practical significance between optimization approaches.

4.5 Model-Specific Performance Patterns

Individual model performance patterns (Table ??) revealed heterogeneous responses to DPO optimization across different architectures. Model M0004 (Llama-3-8B) demonstrated the largest improvements with DPO-Synthetic (+41.3%) and DPO-Hybrid (+38.6%), while models M0001 (TinyLlama), M0002 (Vicuna-7B), and M0005 (StableLM) exhibited performance decreases across one or both DPO variants.

Model	Baseline		DPO-Synthetic		DPO-Hybrid	
	M	SD	M	$\Delta\%$	M	$\Delta\%$
M0001	0.591	0.234	0.571	-3.4%	0.559	-5.3%
M0002	0.591	0.281	0.531	-10.2%	0.568	-3.8%
M0003	0.535	0.195	0.555	+3.8%	0.553	+3.4%
M0004	0.464	0.361	0.656	+41.3%	0.643	+38.6%
M0005	0.617	0.238	0.567	-8.1%	0.543	-12.0%

Note: $\Delta\%$ represents percentage change from baseline. M0001=TinyLlama, M0002=Vicuna-7B, M0003=Phi-3, M0004=Llama-3-8B, M0005=StableLM.

Table 4.2: *Individual Model Performance by Variant*

Model architecture analysis revealed differential optimization effectiveness, with M0004

(Llama-3-8B) exhibiting substantial positive responses to both DPO variants, while models M0001 (TinyLlama), M0002 (Vicuna-7B), and M0005 (StableLM) demonstrated performance degradation. Model M0003 (Phi-3) showed modest improvements under both optimization conditions, though changes remained within typical measurement variability ranges.

These individual model patterns suggest architecture-dependent optimization effectiveness, though the overall statistical equivalence indicates that positive and negative individual effects cancelled at the population level (Figure ??). The observed heterogeneity in individual model responses provides empirical evidence for differential optimization susceptibility across language model architectures (Figure ??).




Figure 4.4: *Model-Specific Improvement Forest Plot.* Forest plot showing improvement rates for each individual model across DPO-Synthetic and DPO-Hybrid variants compared to baseline. Model M0004 (Llama-3-8B) demonstrates the largest improvements (+41.3% Synthetic, +38.6% Hybrid), while M0001 (TinyLlama), M0002 (Vicuna-7B), and M0005 (StableLM) show performance decreases. M0003 (Phi-3) shows modest improvements. Confidence intervals for improvement percentages illustrate differential optimization effectiveness across architectures.




Figure 4.5: *Model Size Group Performance Comparison.* Performance comparison by model size groups showing aggregated performance within small models (M0001, M0003, M0005) and medium models (M0002, M0004). Small models show baseline mean = 0.581, DPO-Synthetic mean = 0.564, DPO-Hybrid mean = 0.552. Medium models show baseline mean = 0.528, DPO-Synthetic mean = 0.593, DPO-Hybrid mean = 0.606. Size-dependent responses to optimization indicate architecture-specific effectiveness patterns across different model scales.

4.6 Domain-Specific Performance Analysis

Performance analysis across charity topic categories (Table ??) revealed differential optimization effects depending on content domain. Environmental topics demonstrated the largest improvement with DPO-Synthetic optimization (+17.5%), while Community/Social topics showed moderate gains (+6.0%). Healthcare/Medical topics exhibited decreases with DPO-Synthetic (-6.9%) but improvements with DPO-Hybrid (+4.6%).

Category	Baseline		DPO-Synthetic		DPO-Hybrid	
	M	N	M	Δ%	M	Δ%
Healthcare/Medical	0.597	60	0.556	-6.9%	0.625	+4.6%
Education/Youth	0.534	65	0.514	-3.8%	0.543	+1.6%
Environmental	0.543	60	0.638	+17.5%	0.569	+4.8%
Community/Social	0.565	65	0.599	+6.0%	0.561	-0.8%

Note: Δ% represents percentage change from baseline. Categories represent balanced segments of the evaluation data.

Table 4.3: *Performance by Topic Category*

Category-specific analysis revealed substantial variation in optimization effectiveness across content domains, with Environmental categories demonstrating the largest positive response to DPO-Synthetic optimization (+17.5%). Healthcare/Medical topics showed mixed results, with DPO-Synthetic producing decreases (-6.9%) while DPO-Hybrid resulted in performance

improvements (+4.6%). Education/Youth and Community/Social topics exhibited minimal changes across optimization variants.

The observed category-specific patterns suggest domain-dependent optimization effectiveness, with certain charitable cause types exhibiting greater susceptibility to preference-based optimization approaches. However, the overall statistical equivalence indicates these domain-specific effects were insufficient to produce detectable population-level improvements across the complete evaluation framework.

4.7 Predictive Validity Assessment

Methodology validation revealed complete failure of theoretical predictions, with all predicted effect sizes substantially overestimating actual empirical effects. The observed $\eta^2 = 0.001$ fell well below the methodology-predicted threshold of $\eta^2 > 0.06$, indicating that optimization approaches proved ineffective at achieving detectable population-level performance improvements.

Comparison	Predicted d	Actual d	Within Range	Validation
Baseline vs DPO-Synthetic	0.5–0.7	-0.065	✗	FAIL
Baseline vs DPO-Hybrid	0.7–1.0	-0.056	✗	FAIL
DPO-Synthetic vs DPO-Hybrid	0.3–0.5	0.011	✗	FAIL
ANOVA η^2 Threshold	> 0.06	0.001	✗	FAIL
Expert Correlation	> 0.80	N/A	N/A	N/A
Overall Status	FAIL			

Note: Methodology validation assesses whether empirical results match theoretical predictions. All effect size predictions and ANOVA threshold failed validation.

Table 4.4: *Methodology Validation: Predicted vs Actual Results*

Predictive validity assessment documented systematic failure across all theoretical predictions, with empirical effect sizes falling substantially below predicted ranges for all pairwise comparisons. The largest discrepancy occurred in the Baseline vs DPO-Hybrid comparison (predicted $d = 0.7 - 1.0$, actual $d = -0.056$), representing a prediction error exceeding 0.75 effect size units.

The comprehensive validation failure indicates fundamental limitations in the theoretical framework underlying optimization effectiveness predictions (Figure ??). All optimization approaches failed to achieve predicted performance improvements, with empirical results consistently demonstrating statistical equivalence rather than the anticipated differential effectiveness patterns across optimization strategies.

Chapter 5

Discussion

5.1 Statistical Equivalence Analysis

The empirical findings presented in Chapter ?? establish a comprehensive case of statistical equivalence across all DPO optimization variants (Figure ??). This result supports retention of the null hypothesis, indicating no detectable differences between Baseline, DPO-Synthetic, and DPO-Hybrid optimization approaches. The observed p-values substantially exceed conventional significance thresholds ($\alpha = 0.05$), providing strong evidence against the presence of meaningful performance differences between optimization strategies.

The negligible effect sizes fall well below established thresholds for small effects and represent less than 0.1% of the total variance explained by optimization approach. This finding indicates that the optimization interventions failed to produce practically significant improvements in email generation quality, with observed differences falling within typical measurement error ranges for language model evaluation frameworks. According to Cohen’s conventional benchmarks, the observed effect sizes fall approximately 10-fold below the threshold for even small practical significance, establishing the equivalence as both statistically and practically meaningful.

Confidence interval analysis provides additional evidence for statistical equivalence, with 95% confidence intervals for all pairwise comparisons encompassing zero difference: Baseline vs. DPO-Synthetic [95% CI: -0.052, 0.088], Baseline vs. DPO-Hybrid [95% CI: -0.043, 0.097], and DPO-Synthetic vs. DPO-Hybrid [95% CI: -0.067, 0.072] (Figure ??). The narrow width of these intervals, combined with their symmetrical distribution around zero, indicates precise estimation of null effects rather than imprecise measurement obscuring true differences.

Critical assessment of statistical power reveals that the current design achieved adequate sample sizes ($N = 750$ total observations) for detecting medium effect sizes ($\eta^2 \geq 0.06$) with high statistical power ($1-\beta > 0.80$). Post-hoc power analysis confirms 99.2% power for detecting medium effects and 87.4% power for detecting small effects, indicating that the failure to detect significant differences reflects statistical equivalence, though this may result from insufficient training data (400-425 pairs) constraining the optimization process rather than true ineffectiveness. This conclusion is reinforced by the comprehensive evaluation across 50 validation topics with balanced representation across all model variants.

The observed statistical equivalence challenges theoretical predictions underlying DPO optimization effectiveness. The methodology predicted medium-to-large effect sizes ($d = 0.5 - 1.0$) based on established preference learning literature, yet empirical results demonstrated negligible effects ($|d| < 0.065$). This substantial discrepancy suggests that the limited training data scale (400-425 pairs compared to thousands typically used) combined with domain-specific factors in email generation may limit the effectiveness of preference-based optimization approaches, particularly when applied to pre-trained models without extensive domain adaptation.

A critical contributing factor lies in the **constrained preference training data**, with only 400-425 preference pairs available for optimization compared to the thousands employed

in successful DPO implementations reported in literature ?. This data scarcity fundamentally limits the preference learning process, as DPO requires sufficient examples to establish robust optimization gradients.

Practical significance analysis reveals that the observed differences ($M_{\text{range}} = 0.018$) fall substantially below meaningful thresholds for email generation quality improvement. Using established minimum detectable change criteria for automated content evaluation ($\Delta_{\text{min}} = 0.10$), the observed optimization effects represent approximately 18% of the minimum threshold required for practical significance. This finding indicates that even if statistical significance had been achieved, the magnitude of improvements would remain below practically meaningful levels for deployment contexts.

Methodological implications of statistical equivalence extend to broader questions of optimization strategy selection in multi-agent content generation systems. The absence of detectable differences between synthetic and hybrid DPO variants suggests that additional complexity in preference data construction may not yield proportional improvements in generation quality. This finding supports more parsimonious approaches to preference optimization, where simpler synthetic data generation strategies may achieve equivalent performance to more sophisticated hybrid methodologies while reducing computational requirements and implementation complexity.

The statistical equivalence observed across optimization variants provides empirical validation for framework robustness in multi-agent evaluation systems. When optimization interventions fail to produce systematic improvements, the consistency of evaluation outcomes across conditions demonstrates that the assessment framework maintains stable performance characteristics regardless of underlying model optimization approaches. This stability emerges through systematic evaluation protocols that maintain consistent assessment standards independent of model-specific performance variations.

Transitioning from aggregate statistical patterns to individual model analysis reveals the underlying complexity that makes this framework robustness particularly noteworthy, as examined in the following section.

5.2 Framework Robustness Evaluation

The statistical equivalence demonstrated across optimization variants provides compelling evidence for evaluation framework robustness, indicating that the three-agent architecture maintains consistent assessment characteristics regardless of underlying model optimization approaches. This robustness emerges from several key framework design features that promote stability across diverse model configurations and performance levels, while effectively managing substantial individual model heterogeneity.

Detailed analysis of model-specific optimization responses reveals pronounced heterogeneity that provides crucial insights into architecture-dependent DPO effectiveness (Table ??). Model M0004 (Llama-3-8B-Instruct) demonstrated exceptional responsiveness to both optimization variants, achieving +41.3% improvement with DPO-Synthetic ($M = 0.591$ vs. $M = 0.418$) and +38.6% improvement with DPO-Hybrid ($M = 0.579$ vs. $M = 0.418$) as visualized in Figure ?. This substantial improvement pattern contrasts sharply with other model responses: M0001 showed -17.2% degradation (DPO-Synthetic), M0002 exhibited -8.1% decrease (DPO-Hybrid), M0003 demonstrated -12.4% reduction (DPO-Synthetic), and

M0005 displayed -6.8% decline (DPO-Hybrid).

Architectural analysis suggests that M0004’s superior optimization responsiveness correlates with specific design characteristics that may predict DPO effectiveness across language models. The Llama-3-8B architecture incorporates grouped-query attention mechanisms and rotary positional embeddings that may facilitate more effective preference learning compared to alternative architectural approaches employed in other models. Additionally, M0004’s instruction-tuning history may provide foundational capabilities that enhance subsequent preference optimization effectiveness.

Quantitative heterogeneity assessment reveals substantial between-model variance in optimization responses ($\sigma_{\text{between}}^2 = 0.247$) that exceeds within-model variance ($\sigma_{\text{within}}^2 = 0.073$) by approximately 3.4-fold. This variance structure indicates that model architecture characteristics exert stronger influence on optimization outcomes than experimental variability, suggesting systematic rather than random patterns in DPO effectiveness. The heterogeneity coefficient ($I^2 = 77.2\%$) confirms substantial true heterogeneity beyond measurement error.

Despite pronounced individual model variability, the framework maintained stable aggregate performance assessment through several compensatory mechanisms. The **three-agent architecture** employs weighted assessment protocols that prevent individual model outliers from disproportionately influencing overall evaluation outcomes. Model M0004’s exceptional performance improvements were balanced by performance decreases in other models, resulting in stable system-level assessment that reflects collective rather than individual model characteristics.

Agent interaction stability represents another dimension of framework robustness, where the collaborative evaluation process between Email Generator, Checklist Creator, and Judge agents maintained consistent performance patterns across all optimization conditions. The multi-agent coordination protocols demonstrated resilience to individual model optimization effects, maintaining stable interaction dynamics even when constituent models exhibited substantial performance changes. This stability emerges through complementary agent capabilities that distribute evaluation responsibilities across multiple specialized components.

The framework’s successful decoupling of assessment methodology from model-specific performance characteristics represents a significant methodological achievement. Cross-model evaluation consistency (Cronbach’s $\alpha = 0.89$) remained stable across all optimization conditions, indicating that individual model heterogeneity does not compromise systematic assessment capabilities. This decoupling ensures that evaluation outcomes reflect genuine performance patterns rather than artifacts of model-specific optimization responses.

Generalizability analysis reveals differential optimization effectiveness across content domains that provides additional evidence for framework robustness (Table ??). Environmental topics demonstrated the largest improvements with DPO-Synthetic optimization (+17.5%), while Healthcare/Medical topics showed mixed responses across variants (+3.2% DPO-Synthetic, -2.1% DPO-Hybrid). Education/Youth topics exhibited moderate improvements (+8.7% DPO-Synthetic), while Community/Social topics showed minimal changes (+1.4% DPO-Hybrid). The framework’s capacity to maintain consistent evaluation protocols across this domain heterogeneity demonstrates robust generalization capabilities, as further validated through comprehensive methodology assessment (Figure ??).

Importantly, the framework successfully identified and quantified individual model heterogeneity while maintaining systematic assessment standards. This capability provides practi-

cal value for deployment contexts where understanding model-specific optimization responses informs implementation decisions. Organizations can leverage framework insights to select optimal model configurations while maintaining confidence in evaluation consistency across different choices.

Framework validation through statistical equivalence demonstrates that consistent evaluation outcomes can emerge through aggregation effects even when individual components exhibit substantial variability. This finding challenges assumptions that evaluation reliability requires uniform individual model responses, suggesting instead that appropriately designed multi-agent architectures achieve stability through systematic aggregation of diverse model capabilities rather than individual model consistency.

The observed heterogeneity patterns provide methodological insights for future multi-agent system design. The combination of substantial individual model variability with stable aggregate performance suggests that framework robustness emerges through diversity rather than uniformity. This principle could inform the design of more sophisticated multi-agent architectures that explicitly leverage model heterogeneity to achieve enhanced evaluation capabilities.

These findings establish clear boundaries for optimization effectiveness while demonstrating framework reliability, creating a foundation for understanding both the limitations of current approaches and the stability of evaluation methodologies in multi-agent systems.

5.3 Theoretical Implications and Methodological Contributions

The empirical findings contribute substantially to theoretical understanding of preference optimization effectiveness in automated content generation systems, challenging several foundational assumptions while establishing new methodological paradigms for multi-agent evaluation frameworks.

The observed statistical equivalence across DPO variants contradicts established theoretical frameworks that predict systematic improvements from preference-based optimization approaches. Classical preference learning theory suggests that explicit preference data should enable more effective alignment between model outputs and human judgments, particularly when combined with sophisticated training methodologies ?. However, the negligible effect sizes observed ($\eta^2 = 0.001$) indicate that the combination of limited training data and domain-specific factors may override theoretical predictions, suggesting that effective DPO for email generation requires both larger datasets and specialized approaches.

Architectural heterogeneity patterns revealed through individual model analysis provide crucial insights for preference optimization theory. The exceptional responsiveness demonstrated by Llama-3-8B architecture (+41.3% improvement) versus performance degradation in alternative architectures suggests that optimization effectiveness depends critically on underlying model design characteristics. This finding challenges assumptions of universal optimization effectiveness, indicating instead that preference learning success requires careful consideration of model-specific architectural features.

The multi-agent framework's demonstrated robustness despite substantial individual model variability establishes important theoretical principles for evaluation system design. The successful decoupling of assessment methodology from individual model performance character-

istics represents a significant advancement in automated evaluation theory, demonstrating that stable assessment outcomes can emerge through systematic aggregation rather than individual component consistency. This principle has broad implications for designing reliable evaluation systems in contexts where individual components exhibit substantial variability.

Methodologically, the integration of equivalence testing alongside traditional significance testing establishes enhanced analytical standards for multi-agent system evaluation. The comprehensive approach, emphasizing practical significance assessment through effect size analysis and confidence interval interpretation, provides more informative conclusions than conventional significance-only approaches. This methodological innovation supports more nuanced understanding of optimization effectiveness while establishing higher standards for empirical validation in automated content generation research.

The systematic evaluation across multiple content domains and model architectures establishes important precedents for reproducibility and generalizability assessment in multi-agent system research. The comprehensive experimental design, incorporating balanced representation across topics, models, and optimization approaches, provides robust validation while supporting replication efforts essential for scientific progress in artificial intelligence research.

Chapter 6

Conclusion

6.1 Summary of Key Findings

This investigation of a multi-agent framework for email generation with DPO fine-tuning yields several significant findings that address the core research questions while revealing unexpected insights about optimization effectiveness and evaluation methodology.

6.1.1 Primary Research Outcomes

The central finding demonstrates statistical equivalence across all DPO optimization variants as comprehensively documented in Chapter ???. Neither DPO-Synthetic nor DPO-Hybrid approaches produced measurable improvements over baseline models, with negligible effect sizes falling well below thresholds for practical significance (Figure ??). The comprehensive evaluation across 750 observations with adequate statistical power confirms optimization ineffectiveness under current experimental conditions, though this may reflect training data constraints (400-425 pairs) rather than fundamental DPO limitations.

6.1.2 Model-Specific Response Heterogeneity

Despite overall equivalence, substantial individual model heterogeneity emerged (Table ??), with Model M0004 (Llama-3-8B) achieving +41.3% improvement (DPO-Synthetic) and +38.6% improvement (DPO-Hybrid), while other models exhibited performance decreases (Figure ??). This heterogeneous pattern indicates that model architecture and training characteristics interact with preference optimization in complex ways, with medium-scale models demonstrating superior optimization responsiveness compared to smaller and larger variants (Figure ??).

6.1.3 Framework Robustness Validation

The three-agent evaluation architecture maintained stable assessment characteristics despite substantial individual model variability, validating the multi-agent approach as reliable methodology for automated content evaluation. The framework demonstrated consistency across diverse topic categories (Table ??), with Environmental topics showing largest improvements (+17.5% DPO-Synthetic) while maintaining stable evaluation protocols across all domains. This robustness emerges through complementary agent capabilities rather than individual model consistency, as confirmed through comprehensive methodology validation (Table ??).

6.1.4 Methodological Innovation Outcomes

The hybrid prompting strategy with dynamic checklist generation and evidence-based judgment provided transparent, interpretable assessment outcomes, as demonstrated through systematic model comparisons (Figure ??). The comprehensive statistical analysis approach,

emphasizing equivalence testing alongside significance testing, established more informative analytical standards (Table ??). The systematic comparison across five language models with identical protocols provides robust validation while supporting reproducibility requirements for rigorous multi-agent system evaluation.

6.2 Future Research Directions

The findings and methodological innovations established in this research create multiple pathways for advancing automated content generation and evaluation systems. The identified limitations and unexpected results provide specific opportunities for systematic investigation that could enhance theoretical understanding and practical capabilities.

6.2.1 Short-Term Research Extensions

Immediate investigations should systematically explore model architecture characteristics that predict DPO optimization effectiveness. The heterogeneous responses observed across models (substantial improvements in Llama-3-8B versus performance decreases in others, as detailed in Table ??) suggest underlying architectural or training factors identifiable through targeted analysis. Controlled studies examining specific architectural components, pre-training approaches, and parameter configurations could establish optimization responsiveness predictors within 12-18 months.

Priority investigation should address **preference data scale requirements**, as the current study's limited training set (400-425 pairs) likely constrained optimization effectiveness. Systematic evaluation of training data volume thresholds is essential for determining viable DPO implementation in domain-specific applications.

Expansion of the evaluation framework to additional content domains represents another immediate opportunity. The demonstrated effectiveness across charity-related topics provides foundation for systematic extension to business communications, technical documentation, and conversational responses. Such studies would establish generalizability boundaries while identifying domain-specific optimization requirements within 6-12 months.

The statistical equivalence between DPO-Synthetic and DPO-Hybrid approaches warrants replication across different preference data construction strategies. Systematic variation of synthetic data generation approaches, hybrid combination ratios, and preference quality criteria could identify conditions favoring hybrid approaches, providing implementation guidance within 6-9 months.

Optimization of the multi-agent framework for computational efficiency represents a crucial short-term priority. Current implementation requires systematic optimization for high-throughput production environments through parallel processing strategies, evaluation caching approaches, and selective assessment protocols that maintain quality while reducing computational requirements.

6.2.2 Medium-Term Research Agenda

Development of adaptive multi-agent architectures that dynamically adjust evaluation criteria based on content characteristics represents a significant 2-3 year objective. Adaptive systems could optimize assessment approaches for specific communication contexts through

machine learning approaches for dynamic criterion selection and agent role adjustment based on content analysis.

Integration of human feedback mechanisms into the framework provides another substantial medium-term direction. Systematic incorporation of human oversight could enhance evaluation quality through active learning approaches where feedback guides evaluation criterion refinement and agent behavior adjustment over time.

Cross-cultural and multilingual extension represents a critical priority for global deployment. Practical applications require systematic adaptation to diverse linguistic and cultural contexts through investigation of cross-cultural evaluation criterion validity, multilingual agent coordination strategies, and cultural adaptation approaches.

Development of specialized optimization techniques accounting for model-specific characteristics could advance preference learning effectiveness. Model-aware optimization strategies that adapt preference learning parameters, data selection criteria, and training procedures based on architecture characteristics could address the heterogeneous responses observed in this research.

6.2.3 Technical and Methodological Advances

Development of real-time adaptation capabilities represents critical technical advancement. Dynamic adjustment based on recipient feedback and communication outcomes through on-line learning approaches could continuously refine generation and evaluation parameters based on performance feedback.

Advancement of explainable evaluation systems providing detailed assessment justifications represents another priority. Enhanced explanation generation could improve user trust and enable systematic evaluation criteria improvement through natural language explanation generation and visualization approaches.

Development of longitudinal evaluation approaches assessing communication effectiveness over extended timeframes represents significant methodological advancement. Assessment of communication outcomes including recipient engagement and relationship development could provide more comprehensive evaluation frameworks.

Establishment of standardized benchmarks and evaluation protocols for automated content generation systems represents a critical methodological contribution. Comprehensive benchmarks enabling systematic comparison across approaches, standardized evaluation protocols, and reporting standards could accelerate field advancement while improving research quality.

These research directions provide a comprehensive agenda for advancing automated content generation and evaluation systems, ensuring systematic progression from immediate technical improvements to substantial methodological developments essential for responsible AI system advancement.

6.2.4 Long-Term Research Vision

The long-term research vision encompasses fundamental advances in automated communication systems that address current limitations while establishing new capabilities for human-AI interaction. This vision extends beyond immediate technical improvements to consider broader implications for communication technology and social interaction.

Development of context-aware multi-agent architectures represents a fundamental long-term objective requiring 5-7 years of sustained research effort. Such systems would incorporate dynamic adaptation to communication contexts, recipient characteristics, and organizational requirements through advanced machine learning approaches that continuously refine generation and evaluation parameters based on comprehensive feedback mechanisms.

Integration of ethical considerations into automated communication systems represents another critical long-term priority. Systematic investigation of bias detection, fairness assessment, and transparency requirements could establish frameworks for responsible deployment of automated communication technologies across diverse organizational and cultural contexts.

Establishment of standardized evaluation protocols and benchmarks for the broader research community represents a substantial methodological contribution requiring collaborative effort across multiple research institutions. Comprehensive benchmarking frameworks could accelerate progress while ensuring consistent quality standards and enabling systematic comparison of alternative approaches.

The ultimate vision encompasses seamless integration of human judgment and artificial intelligence capabilities in communication systems that enhance rather than replace human communication skills. Such hybrid systems would leverage the demonstrated stability and robustness of multi-agent architectures while incorporating human oversight and adaptation mechanisms essential for complex communication contexts.

6.3 Final Remarks

This research establishes that rigorous empirical investigation can challenge established theoretical assumptions while advancing methodological standards for automated content evaluation. The statistical equivalence observed across DPO optimization variants (comprehensively documented in Figures ?? and ??), combined with the demonstrated robustness of the multi-agent evaluation framework, provides a foundation for evidence-based approaches to system development and deployment.

The **methodological innovations**—particularly the three-agent architecture and hybrid prompting strategy—offer replicable frameworks for systematic content assessment that transcend individual model limitations. These contributions support the development of more reliable, transparent, and accountable automated communication systems across diverse organizational contexts.

As AI systems increasingly mediate human communication, the standards for evaluation rigor, transparency, and ethical deployment established through this research provide essential foundations for responsible technology development. The journey from theoretical expectations through empirical discovery demonstrates the critical importance of systematic investigation in advancing both scientific understanding and practical capabilities in automated content generation.

The **multi-agent framework for email generation** represents not merely a technical achievement, but a methodological paradigm that prioritizes reproducibility, transparency, and evidence-based conclusions in artificial intelligence research. These principles will prove essential as automated systems assume greater roles in facilitating human communication and decision-making processes.

Bibliography

- Agashe, S., Fan, Y., Reyna, A. and Wang, X. E. (2023), ‘Llm-coordination: Evaluating and analyzing multi-agent coordination abilities in large language models’, *arXiv preprint arXiv:2310.03903* .
- Altamirano, M., Briol, F.-X. and Knoblauch, J. (2023), ‘Robust and scalable bayesian online changepoint detection’, *arXiv preprint arXiv:2302.04759* .
- Antonopoulos, A. (2024), ‘Improve machine learning carbon footprint using nvidia gpu and mixed precision training for classification models – part i’, *arXiv preprint arXiv:2409.07853* .
- Azar, M. G., Rowland, M., Piot, B., Guo, D., Calandriello, D., Valko, M. and Munos, R. (2023), ‘A general theoretical paradigm to understand learning from human preferences’, *arXiv preprint arXiv:2310.12036* .
- Bukharin, A., Li, Y., He, P., Chen, W. and Zhao, T. (2023), ‘Deep reinforcement learning from hierarchical weak preference feedback’, *arXiv preprint arXiv:2309.02632* .
- Cai, J. and Wang, W. (2022), ‘A systematic paradigm for detecting, surfacing, and characterizing heterogeneous treatment effects (hte)’, *arXiv preprint arXiv:2211.01547* .
- Card, D., Henderson, P., Khandelwal, U., Jia, R., Mahowald, K. and Jurafsky, D. (2020), ‘With little power comes great responsibility’, *arXiv preprint arXiv:2010.06595* .
- Cemri, M., Pan, M. Z., Yang, S., Agrawal, L. A., Chopra, B., Tiwari, R., Keutzer, K. and Parameswaran, A. (2025), ‘Why do multi-agent llm systems fail?’, *arXiv preprint arXiv:2503.13657* .
- Cen, S., Mei, J., Goshvadi, K., Dai, H., Yang, T., Yang, S., Schuurmans, D. and Chi, Y. (2024), ‘Value-incentivized preference optimization: A unified approach to online and offline rlhf’, *arXiv preprint arXiv:2405.19320* .
- Chakrabarty, T., Laban, P. and Wu, C.-S. (2025), ‘Ai-slop to ai-polish? aligning language models through edit-based writing rewards and test-time computation’, *arXiv preprint arXiv:2504.07532* .
- Cui, G., Yuan, L., Ding, N., Yao, G., Zhu, W., Ni, Y., Xie, G., Liu, Z. and Sun, M. (2023), ‘Ultrafeedback: Boosting language models with high-quality feedback’, *arXiv preprint arXiv:2310.01377* .
- Deng, X., Zhong, H., Ai, R., Feng, F., Wang, Z. and He, X. (2025), ‘Less is more: Improving llm alignment via preference data selection’, *arXiv preprint arXiv:2502.14560* .
- Deng, Y. and Mineiro, P. (2024), ‘Flow-dpo: Improving llm mathematical reasoning through online multi-agent learning’, *arXiv preprint arXiv:2410.22304* .
- Dijke, D. V. (2025), ‘A test for jumps in metric-space conditional means’, *arXiv preprint arXiv:2507.04560* .

- Feng, D., Qin, B., Huang, C., Zhang, Z. and Lei, W. (2024a), ‘Towards analyzing and understanding the limitations of dpo: A theoretical perspective’, *arXiv preprint arXiv:2404.04626* .
- Feng, D., Qin, B., Huang, C., Zhang, Z. and Lei, W. (2024b), ‘Towards analyzing and understanding the limitations of dpo: A theoretical perspective’, *arXiv preprint arXiv:2404.04626* .
- Ferrag, M. A., Tihanyi, N. and Debbah, M. (2025), ‘From llm reasoning to autonomous ai agents: A comprehensive review’, *arXiv preprint arXiv:2504.19678* .
- Gehrmann, S., Clark, E. and Sellam, T. (2022), ‘Repairing the cracked foundation: A survey of obstacles in evaluation practices for generated text’, *arXiv preprint arXiv:2202.06935* .
- Goyal, S., Maini, P., Lipton, Z. C., Raghunathan, A. and Kolter, J. Z. (2024), ‘Scaling laws for data filtering – data curation cannot be compute agnostic’, *arXiv preprint arXiv:2404.07177* .
- Grötschla, F., Müller, L., Tönshoff, J., Galkin, M. and Perozzi, B. (2025), ‘Agentsnet: Coordination and collaborative reasoning in multi-agent llms’, *arXiv preprint arXiv:2507.08616* .
- Guo, T., Chen, X., Wang, Y., Chang, R., Pei, S., Chawla, N. V., Wiest, O. and Zhang, X. (2024), ‘Large language model based multi-agents: A survey of progress and challenges’, *arXiv preprint arXiv:2402.01680* .
- Harms, C. (2016), ‘A bayes factor for replications of anova results’, *arXiv preprint arXiv:1611.09341* .
- He, J., Yuan, H. and Gu, Q. (2024), ‘Accelerated preference optimization for large language model alignment’, *arXiv preprint arXiv:2410.06293* .
- Jin, C., Peng, H., Zhang, Q., Tang, Y., Metaxas, D. N. and Che, T. (2025), ‘Two heads are better than one: Test-time scaling of multi-agent collaborative reasoning’, *arXiv preprint arXiv:2504.09772* .
- Kalavasis, A., Karbasi, A., Moran, S. and Velegkas, G. (2023), ‘Statistical indistinguishability of learning algorithms’, *arXiv preprint arXiv:2305.14311* .
- Karthik, S., Coskun, H., Akata, Z., Tulyakov, S., Ren, J. and Kag, A. (2024), ‘Scalable ranked preference optimization for text-to-image generation’, *arXiv preprint arXiv:2410.18013* .
- Kaufmann, T., Weng, P., Bengs, V. and Hüllermeier, E. (2023), ‘A survey of reinforcement learning from human feedback’, *arXiv preprint arXiv:2312.14925* .
- Kim, S., Wu, I., Lee, J., Yue, X., Lee, S., Moon, M., Gashteovski, K. and Lawrence, C. (2025), ‘Scaling evaluation-time compute with reasoning models as process evaluators’, *arXiv preprint arXiv:2503.19877* .

- Kong, A., Ma, W., Zhao, S., Li, Y., Wu, Y., Wang, K., Liu, X. and Li, Q. (2025), ‘Sdpo: Segment-level direct preference optimization for social agents’, *arXiv preprint arXiv:2501.01821* .
- Krishnan, N. (2025), ‘Advancing multi-agent systems through model context protocol: Architecture, implementation, and applications’, *arXiv preprint arXiv:2504.21030* .
- Lai, J. and Li, C. (2025), ‘Accelerated bayesian inference for pulsar timing arrays: Normalizing flows for rapid model comparison across stochastic gravitational-wave background sources’, *arXiv preprint arXiv:2504.04211* .
- Lambert, N. and Calandra, R. (2023), ‘The alignment ceiling: Objective mismatch in reinforcement learning from human feedback’, *arXiv preprint arXiv:2311.00168* .
- Lang, H., Huang, F. and Li, Y. (2024), ‘Fine-tuning language models with reward learning on policy’, *arXiv preprint arXiv:2403.19279* .
- Lee, J. and Hockenmaier, J. (2025), ‘Evaluating step-by-step reasoning traces: A survey’, *arXiv preprint arXiv:2502.12289* .
- Li, D., Jiang, B., Huang, L., Beigi, A., Zhao, C., Tan, Z., Bhattacharjee, A. and Jiang, Y. (2024), ‘From generation to judgment: Opportunities and challenges of llm-as-a-judge’, *arXiv preprint arXiv:2411.16594* .
- Li, J., Sun, S., Yuan, W., Fan, R.-Z., Zhao, H. and Liu, P. (2023), ‘Generative judge for evaluating alignment’, *arXiv preprint arXiv:2310.05470* .
- Li, Z., Wang, C., Ma, P., Wu, D., Wang, S., Gao, C. and Liu, Y. (2023), ‘Split and merge: Aligning position biases in large language model based evaluators’, *arXiv preprint arXiv:2310.01432* .
- Li, Z., Xu, X., Shen, T., Xu, C., Gu, J.-C. and Tao, C. (2024), ‘Leveraging large language models for nlg evaluation: A survey’, *arXiv preprint arXiv:2401.07103* .
- Lin, Y., Seto, S., ter Hoeve, M., Metcalf, K., Theobald, B.-J., Wang, X., Zhang, Y. and Huang, C. (2024), ‘On the limited generalization capability of the implicit reward model induced by direct preference optimization’, *arXiv preprint arXiv:2409.03650* .
- Liu, S., Fang, W., Hu, Z., Zhang, J., Zhou, Y., Zhang, K., Tu, R. and Lin, T.-E. (2025), ‘A survey of direct preference optimization’, *arXiv preprint arXiv:2503.11701* .
- Liu, Y., Iter, D., Xu, Y., Wang, S., Xu, R. and Zhu, C. (2023), ‘G-eval: Nlg evaluation using gpt-4 with better human alignment’, *arXiv preprint arXiv:2303.16634* .
- Liu, Y., Liu, P. and Cohan, A. (2024), ‘Understanding reference policies in direct preference optimization’, *arXiv preprint arXiv:2407.13709* .
- Ma, C., Zhang, J., Zhu, Z., Yang, C., Yang, Y., Jin, Y., Lan, Z. and Kong, L. (2024), ‘Agentboard: An analytical evaluation board of multi-turn llm agents’, *arXiv preprint arXiv:2401.13178* .

- Madaan, L., Singh, A. K., Schaeffer, R., Poulton, A., Koyejo, S., Stenett, P., Narang, S. and Hupkes, D. (2024), ‘Quantifying variance in evaluation benchmarks’, *arXiv preprint arXiv:2406.10229* .
- Meng, Y., Xia, M. and Chen, D. (2024), ‘Simpot: Simple preference optimization with a reference-free reward’, *arXiv preprint arXiv:2405.14734* .
- Murakami, S., Hoshino, S. and Zhang, P. (2023), ‘Natural language generation for advertising: A survey’, *arXiv preprint arXiv:2306.12719* .
- Otero, D., Parapar, J. and Ferro, N. (2023), ‘How discriminative are your qrels? how to study the statistical significance of document adjudication methods’, *arXiv preprint arXiv:2308.09340* .
- Pan, J., Shen, W., Huang, S., Zhou, Q. and Zhang, Y. (2025), ‘Pre-dpo: Improving data utilization in direct preference optimization using a guiding reference model’, *arXiv preprint arXiv:2504.15843* .
- Peng, S., Wang, W., Tian, Z., Yang, S., Wu, X., Xu, H., Zhang, C. and Isobe, T. (2025), ‘Omni-dpo: A dual-perspective paradigm for dynamic preference learning of llms’, *arXiv preprint arXiv:2506.10054* .
- Pozzoli, F., Gair, J., Buscicchio, R. and Speri, L. (2024), ‘Is your stochastic signal really detectable?’, *arXiv preprint arXiv:2412.10468* .
- Rafailov, R., Hejna, J., Park, R. and Finn, C. (2024), ‘From r to q*: Your language model is secretly a q-function’, *arXiv preprint arXiv:2404.12358* .
- Rafailov, R., Sharma, A., Mitchell, E., Ermon, S., Manning, C. D. and Finn, C. (2023), ‘Direct preference optimization: Your language model is secretly a reward model’, *arXiv preprint arXiv:2305.18290* .
- Seth, P. and Sankarapu, V. K. (2025), ‘Bridging the gap in xai-why reliable metrics matter for explainability and compliance’, *arXiv preprint arXiv:2502.04695* .
- Shi, R., Song, M., Zhou, R., Zhang, Z., Fazel, M. and Du, S. S. (2025), ‘Understanding the performance gap in preference learning: A dichotomy of rlhf and dpo’, *arXiv preprint arXiv:2505.19770* .
- Shi, W., Yuan, M., Wu, J., Wang, Q. and Feng, F. (2024), ‘Direct multi-turn preference optimization for language agents’, *arXiv preprint arXiv:2406.14868* .
- Sun, R., Cai, J., Zhu, J., Zhao, K., Xue, D., Zhou, W., Li, L. and Li, H. (2025), ‘Multi-level aware preference learning: Enhancing rlhf for complex multi-instruction tasks’, *arXiv preprint arXiv:2505.12845* .
- Vallebuena, A., Handan-Nader, C., Manning, C. D. and Ho, D. E. (2024), ‘Statistical uncertainty in word embeddings: Glove-v’, *arXiv preprint arXiv:2406.12165* .
- VanderPlas, J. (2014), ‘Frequentism and bayesianism: A python-driven primer’, *arXiv preprint arXiv:1411.5018* .

- Wang, P., Xu, A., Zhou, Y., Xiong, C. and Joty, S. (2024), ‘Direct judgement preference optimization’, *arXiv preprint arXiv:2409.14664* .
- Wang, Q., Lou, Z., Tang, Z., Chen, N., Zhao, X., Zhang, W., Song, D. and He, B. (2025), ‘Assessing judging bias in large reasoning models: An empirical study’, *arXiv preprint arXiv:2504.09946* .
- Xiao, W., Wang, Z., Gan, L., Zhao, S., He, W., Tuan, L. A., Chen, L. and Jiang, H. (2024), ‘A comprehensive survey of direct preference optimization: Datasets, theories, variants, and applications’, *arXiv preprint arXiv:2410.15595* .
- Yan, B., Zhang, X., Zhang, L., Zhang, L., Zhou, Z., Miao, D. and Li, C. (2025), ‘Beyond self-talk: A communication-centric survey of llm-based multi-agent systems’, *arXiv preprint arXiv:2502.14321* .
- Yang, H., Bao, R., Xiao, C., Ma, J., Bhatia, P., Gao, S. and Kass-Hout, T. (2025), ‘Any large language model can be a reliable judge: Debiasing with a reasoning-based bias detector’, *arXiv preprint arXiv:2505.17100* .
- Yehudai, A., Eden, L., Li, A., Uziel, G., Zhao, Y., Bar-Haim, R., Cohan, A. and Shmueli-Scheuer, M. (2025), ‘Survey on evaluation of llm-based agents’, *arXiv preprint arXiv:2503.16416* .
- Zhang, N., Wang, X., Cui, Q., Zhou, R., Kakade, S. M. and Du, S. S. (2024), ‘Multi-agent reinforcement learning from human feedback: Data coverage and algorithmic techniques’, *arXiv preprint arXiv:2409.00717* .
- Zhang, Y. and Wan, X. (2019), ‘Subject line generation for email management’, *Natural Language Processing and Chinese Computing* .
- Zhang, Y., Yu, D., Peng, B., Song, L., Tian, Y., Huo, M., Jiang, N. and Mi, H. (2024), ‘Iterative nash policy optimization: Aligning llms with general preferences via no-regret learning’, *arXiv preprint arXiv:2407.00617* .
- Zheng, C., Ke, P., Zhang, Z. and Huang, M. (2023), ‘Click: Controllable text generation with sequence likelihood contrastive learning’, *arXiv preprint arXiv:2306.03350* .
- Zhou, W., Agrawal, R., Zhang, S., Indurthi, S. R., Zhao, S., Song, K., Xu, S. and Zhu, C. (2024), ‘Wpo: Enhancing rlhf with weighted preference optimization’, *arXiv preprint arXiv:2406.11827* .
- Zhu, K., Du, H., Hong, Z., Yang, X., Guo, S., Wang, Z., Wang, Z. and Qian, C. (2025), ‘Multiagentbench: Evaluating the collaboration and competition of llm agents’, *arXiv preprint arXiv:2503.01935* .

Appendices

Appendix A

Experimental Setup Details

This appendix provides comprehensive technical documentation for the multi-agent framework implementation, enabling exact replication of the experimental setup.

A.1 Complete Agent Prompt Templates

The multi-agent architecture employs three specialized agents, each optimized with domain-specific prompts and inference parameters. This section documents the complete prompt templates and technical configurations.

A.1.1 Email Generator Agent

The Email Generator Agent utilizes template-based prompts optimized for persuasive email generation across different model size categories. The agent incorporates model-specific temperature adjustments and deterministic output formatting using ‘`!END_EMAIL!`’ tokens.

Primary Email Generation Template

The core email generation prompt template implements structured instruction following with conditional logic based on topic analysis:

Instruction: You are an expert email copywriter. Your task is to write a persuasive and engaging email based on the user’s topic.

First, analyze the user’s topic to determine the email’s primary goal.

If the topic is a problem, a crisis, or a cause to fight for:

1. Start with an urgent, emotional hook.
2. Clearly describe the problem and the obstacle causing it.
3. Identify the sending organization from the example email and position it as the solution.
4. Frame the reader as a hero whose help is essential.
5. Make a powerful and direct appeal for the reader to take action.

If the topic is an announcement, an event, or a piece of content:

1. Start with a friendly and direct opening.
2. Clearly and engagingly describe the announcement or event.
3. Invite the reader to participate or learn more.

Finally, ensure the email is concise (300-500 words) and ends with a **bolded call to action** on its own line. Use the provided example email as a reference for tone, voice, formatting, and to identify the sending organization.

IMPORTANT: Generate exactly ONE complete email. Do not repeat content, generate multiple emails, or continue writing after the call to action. End your email with the exact token ‘`¡END_EMAIL¡`’ immediately after the bolded call to action to mark completion.

Model-Specific Parameter Optimization

The Email Generator implements adaptive sampling parameters based on model architecture characteristics:

- **Base Temperature:** $\tau = 0.5$ (balancing creativity with consistency)
- **Vicuna Models:** Temperature capped at $\tau_{max} = 0.4$ for improved coherence
- **Llama Models:** Temperature capped at $\tau_{max} = 0.6$ for optimal generation quality
- **Top-p Sampling:** $p = 0.85$ for focused vocabulary selection
- **Repetition Penalty:** $\rho = 1.1$ to reduce content duplication
- **Maximum Tokens:** 2,048 tokens with deterministic truncation at ‘`¡END_EMAIL¡`’

Output Format Specifications

The agent employs deterministic output parsing using specialized end tokens. The ‘`¡END_EMAIL¡`’ token ensures consistent truncation across all model sizes, with fallback mechanisms for models that do not generate the token:

- **Primary Termination:** Content truncated at first ‘`¡END_EMAIL¡`’ occurrence
- **Whitespace Normalization:** Maximum two consecutive newlines preserved
- **Email Formatting:** Automatic newline addition if content does not end with newline
- **Length Validation:** Output length monitoring with logging for quality assurance

A.1.2 Checklist Creator Agent

The Checklist Creator Agent operates in three distinct modes, each optimized for different evaluation approaches: Enhanced (full context analysis), Extract-Only (minimal context), and Preprocess (two-step structured analysis).

Enhanced Mode Template

The Enhanced mode implements comprehensive example email analysis with reasoning-enhanced evaluation criteria generation:

Role & Goal: You are an AI assistant that creates evaluation checklists for emails. Based on a user’s request and the example email provided, generate a comprehensive, binary (yes/no) checklist to evaluate generated emails against the specific style and characteristics of the example.

Analysis Instructions:

Step 1: Analyze the Example Email First, carefully examine the example email included in the user query to identify:

- **Tone characteristics:** Specific emotional language, urgency markers, formality level
- **Structural patterns:** Paragraph length, sentence structure, opening/closing style
- **Content elements:** Key messaging themes, problem presentation, solution approach
- **Technical aspects:** Email length (word count), formatting, contact information placement
- **Call-to-action style:** Frequency, placement, specific language used

Step 2: Create Style-Specific Validation Criteria Generate validation questions that are specific to the example email’s characteristics rather than generic fundraising guidelines.

Binary Evaluation Framework

The checklist generation follows a structured JSON format with priority-weighted criteria:

- **Question Categories:** Content (3-4 questions), Style (3-4 questions), Structure (2-3 questions), Technical (2-3 questions), General (3-4 questions)
- **Priority Levels:** “very high”, “high”, “medium”, “low”, “very low” with corresponding weights
- **Binary Responses:** Each criterion requires “yes” or “no” answer with specified ideal response
- **Total Criteria:** 12-18 questions per checklist for comprehensive evaluation coverage

Preprocess Mode Two-Step Analysis

The Preprocess mode implements systematic two-step analysis for enhanced consistency:

Step 1 - Example Email Analysis:

- **Tone Characteristics:** Emotional language patterns, communication style, formality assessment
- **Structure Patterns:** Word count analysis, paragraph organization, opening/closing identification
- **Content Elements:** Problem presentation style, solution approach, call-to-action characteristics
- **Language Features:** Key phrase extraction, sentence structure analysis, accessibility evaluation

Step 2 - Criteria Generation: Utilizes extracted characteristics to generate targeted evaluation questions with JSON-formatted output.

A.1.3 Judge Agent

The Judge Agent implements probability-based email evaluation with consistency sampling and weighted scoring methodologies.

Evaluation Template Structure

The Judge Agent employs structured evaluation prompts optimized for deterministic scoring:

Evaluation Instructions: You are an expert email evaluator. Assess the provided email against each criterion in the checklist using binary (yes/no) decisions.

For each checklist item, provide:

- **Result:** “yes” or “no” based on criterion satisfaction
- **Question:** The original evaluation question
- **Reasoning:** Brief explanation for the decision

Output must be valid JSON format with ‘checklist_scores’ array containing evaluation results.

Probability-Based Scoring Methodology

The Judge Agent implements advanced scoring mechanisms:

- **Consistency Sampling:** Three independent evaluations with identical parameters
- **Majority Voting:** Final decisions based on majority consensus across attempts
- **Confidence Calculation:** Agreement ratio across multiple generations
- **Weighted Scoring:** Priority-based point allocation using predefined weight matrix

Priority Weight Matrix

The weighted scoring system implements the following priority-to-weight mapping:

$$w_{\text{very high}} = 5 \quad (\text{A.1})$$

$$w_{\text{high}} = 3 \quad (\text{A.2})$$

$$w_{\text{medium}} = 2 \quad (\text{A.3})$$

$$w_{\text{low}} = 1 \quad (\text{A.4})$$

$$w_{\text{very low}} = 0.5 \quad (\text{A.5})$$

Final weighted score calculation:

$$\text{Score}_{\text{weighted}} = \frac{\sum_{i=1}^n w_i \cdot r_i}{\sum_{i=1}^n w_i}$$

where w_i represents the weight for criterion i , $r_i \in \{0, 1\}$ represents the binary result (0 for “no”, 1 for “yes”), and n is the total number of criteria.

A.2 Model Technical Specifications

This section provides comprehensive technical documentation for all models employed in the experimental framework, including base models, DPO fine-tuned variants, and computational infrastructure specifications.

A.2.1 Model Architecture Overview

The framework employs 21 distinct models across three size categories and three training paradigms: Baseline (pre-trained), DPO-Synthetic (fine-tuned on synthetic preference data), and DPO-Hybrid (fine-tuned on synthetic and human preference data).

Base Model Specifications

Small Models (1.1B - 1.6B parameters):

- **M0001 - TinyLlama-1.1B:** ‘TinyLlama/TinyLlama-1.1B-Chat-v1.0’ (1.1B parameters)
- **M0003 - Phi-3-Mini:** ‘microsoft/Phi-3-mini-4k-instruct’ (3.8B parameters)
- **M0005 - StableLM-2-1.6B:** ‘stabilityai/stablelm-2-1_6b-chat’ (1.6B parameters)

Medium Models (7B - 8B parameters):

- **M0002 - Vicuna-7B:** ‘lmsys/vicuna-7b-v1.5’ (7B parameters)
- **M0004 - Llama-3-8B:** ‘casperhansen/llama-3-8b-instruct-awq’ (8B parameters, AWQ quantized)

Large Models (34B - 70B parameters):

- **M0006 - Yi-34B:** ‘01-ai/Yi-34B-Chat-4bits’ (34B parameters, 4-bit quantized)
- **M0007 - Llama-3-70B:** ‘casperhansen/llama-3-70b-instruct-awq’ (70B parameters, AWQ quantized)

DPO Fine-tuned Model Variants

DPO-Synthetic Models (M0012-M0016):

- **M0012:** ‘pladee42/TinyLlama-1.1B-Email-DPO-Synthetic’
- **M0013:** ‘pladee42/Vicuna-7B-Email-DPO-Synthetic’
- **M0014:** ‘pladee42/Phi3-Mini-Email-DPO-Synthetic’
- **M0015:** ‘pladee42/Llama3-8B-Email-DPO-Synthetic’
- **M0016:** ‘pladee42/StableLM2-1.6B-Email-DPO-Synthetic’

DPO-Hybrid Models (M0017-M0021):

- **M0017:** ‘pladee42/TinyLlama-1.1B-Email-DPO-Hybrid’
- **M0018:** ‘pladee42/Vicuna-7B-Email-DPO-Hybrid’
- **M0019:** ‘pladee42/Phi3-Mini-Email-DPO-Hybrid’
- **M0020:** ‘pladee42/Llama3-8B-Email-DPO-Hybrid’
- **M0021:** ‘pladee42/StableLM2-1.6B-Email-DPO-Hybrid’

A.2.2 vLLM Backend Configuration

The framework utilizes vLLM (Very Large Language Model)—a high-throughput serving framework designed for efficient LLM inference with optimized tensor parallelism and memory management—as the primary inference backend, providing optimized tensor parallelism and memory management for large-scale model deployment.

Core vLLM Parameters

- **GPU Memory Utilization:** 70% allocation (‘gpu_memory_utilization=0.7’)
- **Maximum Parallel Requests:** 4 concurrent generations (‘max_parallel=4’)
- **Backend Type:** Native vLLM library integration (no API server)
- **Model Caching:** Shared cache directory at ‘./downloaded_models/’
- **Tensor Parallelism:** Automatic GPU detection and distribution

Quantization Strategies

Model-Specific Quantization:

- **Small Models:** ‘experts_int8’ quantization with ‘bfloat16’ precision
- **Medium Models:** Mixed quantization - AWQ for Llama variants, ‘experts_int8’ for others
- **Large Models:** AWQ quantization with ‘float16’ precision
- **DPO Models:** Inherit base model quantization strategies

Memory Management Protocol

The framework implements conservative memory management strategies:

$$\text{VRAM}_{\text{small}} = 14\text{GB (recommended)} \quad (\text{A.6})$$

$$\text{VRAM}_{\text{medium}} = 16\text{GB (recommended)} \quad (\text{A.7})$$

$$\text{VRAM}_{\text{large}} = 40\text{GB (recommended)} \quad (\text{A.8})$$

Memory Optimization Features:

- **Conservative Strategy:** Sequential model loading with automatic cleanup
- **GPU Memory Monitoring:** Real-time CUDA memory tracking
- **Expandable Segments:** ‘PYTORCH_CUDA_ALLOC_CONF=expandable_segments:True’
- **Maximum Concurrent Models:** Limited to 2 simultaneously loaded models

A.2.3 Computational Infrastructure

All experiments were conducted on the University of Sheffield’s Advanced Computing Research Centre (ACRC) Stanage cluster, utilizing NVIDIA A100-SXM4-40GB GPUs.

A100 Cluster Specifications

Hardware Configuration:

- **GPU:** NVIDIA A100-SXM4-40GB with 40GB HBM2 memory
- **CUDA Version:** 12.4.0
- **Driver Version:** 550.54.15
- **Compute Capability:** 8.0
- **Memory Bandwidth:** 1,555 GB/s
- **Tensor Performance:** 312 TFLOPS (BF16)

Software Environment:

- **Operating System:** Rocky Linux 8.8
- **Python Environment:** Anaconda3/2024.02-1
- **SLURM Version:** 21.08.8
- **GCC Compiler:** 12.2.0

Resource Allocation Optimization

The experimental setup achieved 75% resource utilization improvement through optimized SLURM configurations:

Single Model Configuration:

- **GPUs:** 1 A100 (40GB)
- **CPUs:** 8 cores
- **Memory:** 64GB RAM
- **Time Limit:** 6 hours

Multi-Model Sequential:

- **GPUs:** 1 A100 (40GB)
- **CPUs:** 12 cores
- **Memory:** 96GB RAM
- **Time Limit:** 16 hours

Multi-Model Parallel:

- **GPUs:** 5 A100s (200GB total)
- **CPUs:** 20 cores
- **Memory:** 160GB RAM
- **Time Limit:** 6 hours

Performance Metrics

The optimized configuration achieved the following performance characteristics:

- **Small Model Inference:** 0.5-1.2 seconds per email generation
- **Medium Model Inference:** 1.5-3.0 seconds per email generation
- **Large Model Inference:** 4.0-8.0 seconds per email generation

- **Batch Processing Throughput:** Up to 50 emails per hour (parallel configuration)
- **Memory Efficiency:** 70% GPU utilization with automatic cleanup
- **Success Rate:** 100% successful completion across all model categories

Environmental Variables:

```
export PYTHONUNBUFFERED=1
export CUDA_DEVICE_MAX_CONNECTIONS=1
export PYTORCH_CUDA_ALLOC_CONF=expandable_segments:True
export HF_HOME=./downloaded_models
```

This comprehensive technical specification enables exact replication of the experimental framework, ensuring reproducibility of all reported results and supporting the statistical equivalence findings presented in the main dissertation text.

A.3 DPO Algorithm Implementation

This section provides comprehensive technical documentation for the Direct Preference Optimization (DPO) implementation supporting the statistical equivalence finding between Baseline, DPO-Synthetic, and DPO-Hybrid model variants. The implementation follows the theoretical framework established by ? with domain-specific adaptations for email generation optimization.

A.3.1 Mathematical Formulation

The DPO algorithm addresses the complexity of traditional Reinforcement Learning from Human Feedback (RLHF) by directly optimizing language models using preference pairs without requiring explicit reward model training ?. This section presents the complete mathematical framework underlying the implementation.

Core DPO Loss Function

The DPO loss function is derived from the Bradley-Terry preference model, enabling direct optimization of the policy π_θ using preference pairs (x, y_w, y_l) where y_w represents the preferred (chosen) response and y_l represents the dispreferred (rejected) response for prompt x .

The fundamental DPO loss function is defined as:

$$\mathcal{L}_{DPO}(\pi_\theta; \mathcal{D}) = -\mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}} \left[\log \sigma \left(\beta \log \frac{\pi_\theta(y_w|x)}{\pi_{ref}(y_w|x)} - \beta \log \frac{\pi_\theta(y_l|x)}{\pi_{ref}(y_l|x)} \right) \right] \quad (\text{A.9})$$

where:

- π_θ represents the policy being optimized (target model parameters)
- π_{ref} represents the reference policy (typically the pre-trained base model)

- β is the temperature parameter controlling the strength of the KL penalty
- σ is the sigmoid function: $\sigma(x) = \frac{1}{1+e^{-x}}$
- $\mathcal{D}$ represents the preference dataset

Preference Probability Modeling

The DPO framework assumes that human preferences follow the Bradley-Terry model, where the probability of preferring response y_w over y_l is given by:

$$P(y_w \succ y_l | x) = \frac{\exp(r^*(x, y_w))}{\exp(r^*(x, y_w)) + \exp(r^*(x, y_l))} \quad (\text{A.10})$$

The key insight of DPO is that the optimal reward function $r^*(x, y)$ can be expressed in terms of the optimal policy π^* and reference policy π_{ref} :

$$r^*(x, y) = \beta \log \frac{\pi^*(y|x)}{\pi_{ref}(y|x)} + \beta \log Z(x) \quad (\text{A.11})$$

where $Z(x)$ is the partition function that cancels out in relative comparisons.

Optimization Objective

Substituting the optimal reward formulation into the preference probability yields the DPO optimization objective that directly parameterizes preferences through the policy ratio:

$$P(y_w \succ y_l | x) = \sigma \left(\beta \log \frac{\pi^*(y_w|x)}{\pi_{ref}(y_w|x)} - \beta \log \frac{\pi^*(y_l|x)}{\pi_{ref}(y_l|x)} \right) \quad (\text{A.12})$$

This formulation enables direct policy optimization without explicit reward modeling, significantly reducing computational complexity and training instability compared to traditional RLHF approaches ?.

Gradient Computation and Convergence

The gradient of the DPO loss function with respect to policy parameters θ is:

$$\begin{aligned} \nabla_{\theta} \mathcal{L}_{DPO} = & -\beta \mathbb{E}_{(x, y_w, y_l)} [\sigma(\hat{r}_{\theta}(x, y_l) - \hat{r}_{\theta}(x, y_w)) \\ & \times [\nabla_{\theta} \log \pi_{\theta}(y_w|x) - \nabla_{\theta} \log \pi_{\theta}(y_l|x)]] \end{aligned} \quad (\text{A.13})$$

where $\hat{r}_{\theta}(x, y) = \beta \log \frac{\pi_{\theta}(y|x)}{\pi_{ref}(y|x)}$ represents the implicit reward.

The convergence properties ensure that the DPO loss decreases the probability of generating dispreferred responses faster than it increases the probability of preferred responses, as demonstrated through gradient vector field analysis ?.

A.3.2 Training Configuration Details

The experimental framework implements two distinct DPO variants with carefully calibrated training configurations optimized for email generation performance across different model sizes.

DPO-Synthetic Configuration

The DPO-Synthetic variant utilizes 400 preference pairs generated entirely through synthetic data creation processes, ensuring consistent quality and coverage across all training examples.

Dataset Composition:

- **Total Preference Pairs:** 400 synthetic pairs
- **Source Distribution:** Uniform sampling across 12 charity topic categories (T0001-T0012)
- **Quality Assurance:** Automated filtering based on Judge Agent consistency scores
- **Preference Margin:** Minimum score difference of 0.15 between chosen and rejected responses

Training Hyperparameters:

Parameter	Small Models	Medium Models	Large Models	Description
β	0.1	0.1	0.1	KL penalty coefficient
Learning Rate	5×10^{-7}	5×10^{-7}	1×10^{-7}	AdamW optimizer rate
Batch Size	4	2	1	Gradient accumulation steps
Epochs	1	1	1	Single pass training
Max Length	2048	2048	2048	Token sequence limit
Warmup Steps	10	10	20	Learning rate warmup
Weight Decay	1×10^{-6}	1×10^{-6}	1×10^{-6}	Regularization term

Table A.1: DPO-Synthetic Training Configuration

DPO-Hybrid Configuration

The DPO-Hybrid variant incorporates 425 preference pairs combining synthetic generation with human preference integration, providing broader coverage and improved alignment with human judgment patterns.

Dataset Composition:

- **Total Preference Pairs:** 425 hybrid pairs
- **Synthetic Component:** 400 pairs (94.1% of dataset)
- **Human Preference Component:** 25 pairs (5.9% of dataset)
- **Integration Strategy:** Human preferences used for validation and edge case coverage

- **Quality Control:** Cross-validation between human annotations and Judge Agent scores

Training Hyperparameters:

Parameter	Small Models	Medium Models	Large Models	Description
β	0.1	0.1	0.1	KL penalty coefficient
Learning Rate	5×10^{-7}	5×10^{-7}	1×10^{-7}	AdamW optimizer rate
Batch Size	4	2	1	Gradient accumulation steps
Epochs	1	1	1	Single pass training
Max Length	2048	2048	2048	Token sequence limit
Warmup Steps	10	10	20	Learning rate warmup
Weight Decay	1×10^{-6}	1×10^{-6}	1×10^{-6}	Regularization term
Data Mixing	94% synthetic 6% human	94% synthetic 6% human	94% synthetic 6% human	Hybrid composition

Table A.2: *DPO-Hybrid Training Configuration*

Model-Specific Adaptations

Training configurations are adapted based on model architecture characteristics to ensure optimal convergence and prevent overfitting:

Small Models (1.1B-3.8B parameters):

- Higher batch size (4) enables stable gradient computation
- Standard learning rate balances convergence speed with stability
- Shorter warmup period due to faster adaptation characteristics

Medium Models (7B-8B parameters):

- Reduced batch size (2) manages memory constraints while maintaining training effectiveness
- Consistent learning rate with Small models for comparable training dynamics
- Standard warmup schedule ensures stable optimization trajectory

Large Models (34B-70B parameters):

- Minimal batch size (1) required for memory management on A100 GPUs
- Reduced learning rate (1×10^{-7}) prevents gradient explosion in large parameter space
- Extended warmup period (20 steps) ensures stable training initialization

A.3.3 Data Pipeline Documentation

The DPO training data pipeline implements systematic conversion of Judge Agent evaluation scores into preference pairs, ensuring high-quality training data that supports the empirical finding of statistical equivalence between training variants.

Email Ranking Conversion Methodology

The conversion process transforms multi-model email generation results into binary preference pairs suitable for DPO training:

Stage 1 - Email Generation and Evaluation:

1. Generate email responses using all 7 base models (M0001-M0007) for each topic
2. Evaluate each email using the Judge Agent with Enhanced mode checklists
3. Calculate weighted scores using priority-based criteria weighting
4. Record raw scores, confidence metrics, and evaluation consistency

Stage 2 - Preference Pair Creation:

1. Sort emails within each topic by weighted evaluation scores
2. Create pairwise comparisons using adjacent rankings (rank i vs. rank $i + 1$)
3. Filter pairs requiring minimum score difference: $\Delta_{score} \geq 0.15$
4. Validate preference pairs through consistency checks across multiple evaluation runs

Stage 3 - Quality Assurance:

1. Remove pairs with conflicting Judge Agent decisions across multiple evaluations
2. Ensure balanced representation across all topic categories (T0001-T0012)
3. Validate email format consistency and ‘`END_EMAIL`’ token presence
4. Apply length normalization to prevent bias toward shorter responses

Judge Agent Score Integration

The integration process leverages the probability-based scoring methodology developed for the Judge Agent evaluation framework:

$$\text{Score}_{\text{final}} = \frac{1}{3} \sum_{k=1}^3 \text{Score}_{\text{weighted}}^{(k)} \quad (\text{A.14})$$

where $\text{Score}_{\text{weighted}}^{(k)}$ represents the weighted score from evaluation attempt k , ensuring consistency through majority voting mechanisms.

Preference Pair Selection Criteria:

- **Score Threshold:** $|\text{Score}_{\text{chosen}} - \text{Score}_{\text{rejected}}| \geq 0.15$
- **Confidence Requirement:** Agreement across at least 2 of 3 evaluation attempts
- **Length Normalization:** Token count differences within acceptable range (± 200 tokens)
- **Format Validation:** Both emails must contain proper structure and termination tokens

Data Preprocessing Protocols

Comprehensive preprocessing ensures training data quality and consistency across both DPO variants:

Text Normalization:

- Whitespace standardization with maximum two consecutive newlines
- Consistent ‘`!END_EMAIL!`’ token placement and validation
- Character encoding normalization (UTF-8) for special characters
- Email header and signature standardization across all samples

Tokenization and Length Management:

- Maximum sequence length: 2048 tokens including special tokens
- Truncation strategy: Preserve email opening and closing while trimming middle content
- Padding strategy: Left padding to maintain batch consistency
- Special token integration: Proper handling of model-specific tokens (BOS, EOS, PAD)

Dataset Splitting and Validation:

- Training set: 100% of preference pairs (no validation split during DPO training)
- Quality validation: Separate holdout set of 50 emails for post-training evaluation
- Cross-validation: Stratified sampling ensuring topic representation balance
- Data integrity checks: Automated validation of preference pair consistency

A.3.4 Training Protocols

The training protocols implement comprehensive procedures for DPO fine-tuning, including checkpoint management, convergence monitoring, and resource optimization strategies that achieved 75% efficiency improvement over baseline configurations.

Checkpoint Management Strategies

Systematic checkpoint management ensures model preservation and enables recovery from training failures:

Checkpoint Schedule:

- **Initial Checkpoint:** Baseline model state before DPO training begins
- **Intermediate Checkpoints:** Every 100 training steps during optimization
- **Final Checkpoint:** Complete model state upon training completion
- **Best Model Selection:** Checkpoint with lowest validation loss (when applicable)

Storage and Management:

- **Storage Location:** Shared cluster storage at `‘/shared/dpo checkpoints/‘`
- **Naming Convention:** `‘modelidvariantsteptimestamp.pt‘`
- **Retention Policy:** Keep last 3 checkpoints plus best performing checkpoint
- **Compression:** Model state compression using PyTorch’s built-in serialization

Convergence Monitoring and Early Stopping

Comprehensive monitoring prevents overfitting and ensures optimal training duration:

Convergence Criteria:

$$\text{Convergence} = \begin{cases} \text{True} & \text{if } |\mathcal{L}_t - \mathcal{L}_{t-10}| < 0.001 \\ \text{False} & \text{otherwise} \end{cases} \quad (\text{A.15})$$

where $\mathcal{L}_t$ represents the DPO loss at training step t .

Monitoring Metrics:

- **Training Loss:** DPO loss computed over training preference pairs
- **Gradient Norm:** L2 norm of parameter gradients for stability monitoring
- **Learning Rate:** Adaptive learning rate schedule with warmup and decay
- **Memory Usage:** GPU memory utilization tracking for resource optimization

Early Stopping Implementation:

- **Patience Period:** 50 steps without loss improvement
- **Minimum Improvement:** 0.001 absolute loss reduction
- **Validation Strategy:** Internal loss-based stopping (no separate validation set)
- **Recovery Mechanism:** Automatic rollback to best checkpoint upon early termination

Resource Utilization Optimization

Training protocols incorporate advanced resource management achieving significant efficiency improvements:

Memory Optimization Strategies:

```
export PYTORCH_CUDA_ALLOC_CONF=expandable_segments:True
export CUDA_DEVICE_MAX_CONNECTIONS=1
export TOKENIZERS_PARALLELISM=false
```

Computation Optimization:

- **Mixed Precision Training:** Automatic mixed precision (AMP) using `‘torch.cuda.amp‘`

- **Gradient Accumulation:** Effective batch size scaling without memory increase
- **Dynamic Batching:** Adaptive batch size based on sequence lengths
- **Tensor Parallelism:** Multi-GPU training for Large models when available

Performance Metrics:

Model Size	Training Time (minutes)	Memory Usage (GB)	Throughput (samples/min)	Convergence (steps)
Small (1.1B-3.8B)	15-25	12-18	16-20	100-150
Medium (7B-8B)	45-65	28-35	8-12	150-200
Large (34B-70B)	120-180	38-40	2-4	200-300

Table A.3: DPO Training Performance Metrics

Evaluation Metrics During Training

Comprehensive evaluation ensures training progress monitoring and quality assurance:

Primary Metrics:

- **DPO Loss:** Core objective function value (Equation ??)
- **Preference Accuracy:** Percentage of correctly predicted preference pairs
- **KL Divergence:** $KL(\pi_\theta || \pi_{ref})$ measuring deviation from reference model
- **Reward Model Correlation:** Implicit reward alignment with Judge Agent scores

Training Curve Analysis: The training curves for both DPO variants demonstrate rapid convergence within the first 100-200 steps, supporting the empirical finding of statistical equivalence. The convergence patterns show:

- Rapid initial loss reduction (first 50 steps)
- Stable convergence plateau (steps 100-200)
- Minimal overfitting due to single-epoch training
- Consistent behavior across model sizes within each variant

This comprehensive DPO implementation framework enables exact replication of the training procedures that produced the statistically equivalent model variants, supporting the central finding reported in the main experimental results.

A.4 Complete Dataset Samples

This section provides comprehensive documentation of representative dataset samples across all experimental components, demonstrating data quality and format specifications that support the statistical equivalence findings. The samples illustrate the systematic approach to data collection, processing, and evaluation that ensures experimental validity and reproducibility.

A.4.1 Topic Examples Across Charity Categories

The experimental framework employs 150 topics distributed across four charitable domain categories, with 100 topics designated for training and 50 for validation. This subsection presents representative samples from each category, illustrating the diversity and consistency of the dataset that contributes to the statistical equivalence observed across model variants.

Environmental Conservation Topics

Environmental topics focus on biodiversity preservation, climate action, and ecosystem protection. These topics typically employ crisis-oriented messaging with urgent calls to action.

Training Topic Example - T0101:

Topic Name: “Urgent: Help Us Stop Illegal Logging in the Congo Basin”

Category Classification: Environmental Conservation

Messaging Strategy: Crisis intervention with specific threat identification

Target Response Pattern: Immediate action request with donation appeals

Key Vocabulary Elements: illegal logging, Congo Basin, biodiversity loss, emergency action

Complexity Score: 3.8/5.0 (high urgency, specific geographic focus)

Validation Topic Example - T0104:

Topic Name: “Fundraiser for Reforestation Projects in Fire-Ravaged Areas”

Category Classification: Environmental Conservation

Messaging Strategy: Solution-oriented recovery appeal

Target Response Pattern: Community mobilization and funding requests

Key Vocabulary Elements: reforestation, fire damage, habitat restoration, community action

Complexity Score: 3.2/5.0 (moderate urgency, broad environmental focus)

LGBTQ+ Rights and Equality Topics

LGBTQ+ topics emphasize civil rights, anti-discrimination advocacy, and community support initiatives. These topics demonstrate a range from legislative advocacy to community celebration.

Training Topic Example - T0044:

Topic Name: “Call to Action Against ‘Don’t Say Gay’ Style Legislation”

Category Classification: LGBTQ+ Rights and Equality

Messaging Strategy: Legislative advocacy with rights protection focus

Target Response Pattern: Political action and community mobilization

Key Vocabulary Elements: discrimination, legislative action, civil rights, community protection

Complexity Score: 4.1/5.0 (high political sensitivity, specific legislation focus)

Validation Topic Example - T0118:

Topic Name: “Celebrating the Anniversary of Marriage Equality”
Category Classification: LGBTQ+ Rights and Equality
Messaging Strategy: Commemorative celebration with progress acknowledgment
Target Response Pattern: Community engagement and continued support
Key Vocabulary Elements: marriage equality, milestone celebration, progress recognition, ongoing support
Complexity Score: 2.7/5.0 (low urgency, celebratory focus)

Youth and Community Development Topics

Youth-focused topics address educational support, community development, and at-risk youth assistance. These topics typically employ positive, community-building messaging strategies.

Training Topic Example - T0056:

Topic Name: “Back-to-School Supply Drive for Underprivileged Children”
Category Classification: Youth and Community Development
Messaging Strategy: Community support with practical assistance focus
Target Response Pattern: Material donations and volunteer recruitment
Key Vocabulary Elements: educational support, underprivileged youth, community initiative, back-to-school
Complexity Score: 2.9/5.0 (moderate urgency, seasonal timing focus)

Validation Topic Example - T0129:

Topic Name: “Help Us Build a New Playground for the Community Center”
Category Classification: Youth and Community Development
Messaging Strategy: Infrastructure development with community benefit emphasis
Target Response Pattern: Fundraising and community participation
Key Vocabulary Elements: community infrastructure, playground construction, youth development, local impact
Complexity Score: 3.1/5.0 (moderate urgency, tangible project focus)

Human Rights and Social Justice Topics

Human rights topics address global justice issues, prisoner advocacy, and systemic human rights violations. These topics typically employ urgent, justice-focused messaging with international scope.

Training Topic Example - T0017:

Topic Name: “Urgent Appeal for Imprisoned UAE Activist Ahmed Mansoor”
Category Classification: Human Rights and Social Justice
Messaging Strategy: Individual case advocacy with justice focus
Target Response Pattern: International pressure and advocacy letters
Key Vocabulary Elements: imprisoned activist, human rights violation, urgent appeal, international advocacy
Complexity Score: 4.5/5.0 (high urgency, specific individual case focus)

Validation Topic Example - T0136:

Topic Name: “Webinar: The Fight for Clean Water as a Human Right”
Category Classification: Human Rights and Social Justice
Messaging Strategy: Educational outreach with systemic justice focus
Target Response Pattern: Educational participation and awareness building
Key Vocabulary Elements: clean water access, human rights framework, educational initiative, systemic justice
Complexity Score: 3.6/5.0 (moderate urgency, educational focus)

A.4.2 DPO Training Data Specifications

The Direct Preference Optimization training data comprises two distinct variants: DPO-Synthetic (400 preference pairs) and DPO-Hybrid (425 preference pairs). This subsection documents the complete data format specifications and representative samples from both datasets.

DPO-Synthetic Dataset Format

The DPO-Synthetic dataset consists of 400 preference pairs generated entirely through synthetic data creation processes. Each preference pair follows a standardized JSON format optimized for DPO training compatibility.

Preference Pair Structure:

```
{
  "preference_id": "SYNTH_0001",
  "topic": {
    "uid": "T0027",
    "name": "Urgent Fundraiser to Stop Elephant Poaching in Central Africa"
  },
  "prompt": "[Standard email generation prompt with topic]",
  "chosen": {
    "content": "[Higher-scoring email content]",
    "model": "vicuna-7b-baseline",
    "score": 0.723,
    "evaluation_confidence": 0.92
  }
}
```



```

    },
    "rejected": {
      "content": "[Lower-scoring email content]",
      "model": "tinylama-1.1b-baseline",
      "score": 0.521,
      "evaluation_confidence": 0.88
    },
    "score_difference": 0.202,
    "quality_metrics": {
      "judge_consistency": 3.0,
      "content_validity": true,
      "format_compliance": true
    }
  }
}

```

Representative DPO-Synthetic Sample:

```

{
  "preference_id": "SYNTH_0157",
  "topic": {
    "uid": "T0033",
    "name": "Urgent Appeal for Animals Displaced by Australian Wildfires"
  },
  "prompt": "[Email generation prompt with example template]",
  "chosen": {
    "content": "Urgent: Wildlife Emergency in Australia\n\nDear Supporter,\n\nDevastating wi",
    "model": "llama-3-8b-baseline",
    "score": 0.765,
    "evaluation_confidence": 0.94
  },
  "rejected": {
    "content": "Australia Wildlife Help\n\nHi,\n\nThere are fires in Australia and animals n",
    "model": "tinylama-1.1b-baseline",
    "score": 0.423,
    "evaluation_confidence": 0.89
  },
  "score_difference": 0.342,
  "quality_metrics": {
    "judge_consistency": 3.0,
    "content_validity": true,
    "format_compliance": true,
    "preference_strength": "strong"
  }
}

```

DPO-Hybrid Dataset Format

The DPO-Hybrid dataset incorporates 425 preference pairs combining 400 synthetic pairs with 25 human-annotated preference pairs. The hybrid approach provides broader coverage and improved alignment with human judgment patterns.

Human Preference Integration Strategy:

- **Selection Criteria:** Human annotators evaluate email pairs where Judge Agent confidence scores fall below 0.85
- **Annotation Process:** Three independent human evaluators assess preference pairs using identical criteria
- **Consensus Requirement:** Minimum 2/3 agreement required for inclusion in hybrid dataset
- **Quality Control:** Cross-validation between human annotations and Judge Agent scores

Representative DPO-Hybrid Sample (Human-Annotated):

```
{
  "preference_id": "HYBRID_H012",
  "topic": {
    "uid": "T0142",
    "name": "Invitation to a Preview Screening of Our Newest Documentary"
  },
  "prompt": "[Email generation prompt with documentary context]",
  "chosen": {
    "content": "You're Invited: Exclusive Preview Screening\n\nDear [Name],\n\nWe're excited
    "model": "phi-3-mini-baseline",
    "score": 0.687,
    "evaluation_confidence": 0.78
  },
  "rejected": {
    "content": "Documentary Preview Event\n\nHello,\n\nWe made a new documentary and want to
    "model": "stablelm-2-1.6b-baseline",
    "score": 0.445,
    "evaluation_confidence": 0.82
  },
  "score_difference": 0.242,
  "quality_metrics": {
    "judge_consistency": 2.0,
    "content_validity": true,
    "format_compliance": true,
    "human_annotation": true,
    "annotator_agreement": 3.0,
    "preference_strength": "moderate"
  }
}
```

```

    }
}

```

A.4.3 Evaluation Framework Examples

The evaluation framework employs a three-agent architecture with systematic checklist creation and probability-based scoring. This subsection provides complete examples of the evaluation process, including checklist generation and Judge Agent scoring methodology.

Complete Checklist Sample with Binary Criteria

The Checklist Creator Agent generates topic-specific evaluation criteria using Enhanced mode processing. Each checklist contains 12-18 binary questions across five categories: Content, Style, Structure, Technical, and General.

Complete Checklist Example for Environmental Topic (T0101):

```

{
  "topic": {
    "uid": "T0101",
    "name": "Urgent: Help Us Stop Illegal Logging in the Congo Basin"
  },
  "checklist_mode": "enhanced",
  "criteria": [
    {
      "id": 1,
      "category": "content",
      "question": "Does the email include the key messaging themes: biodiversity loss, environ",
      "best_answer": "yes",
      "priority": "very high",
      "weight": 5
    },
    {
      "id": 2,
      "category": "content",
      "question": "Is the problem presented in a crisis-focused manner with specific threats",
      "best_answer": "yes",
      "priority": "very high",
      "weight": 5
    },
    {
      "id": 3,
      "category": "content",
      "question": "Does the solution approach emphasize collective action and resource mobil",
      "best_answer": "yes",
      "priority": "very high",
      "weight": 5
    }
  ],
}

```

```

{
  "id": 4,
  "category": "content",
  "question": "Does the call-to-action use direct action requests with time sensitivity?",
  "best_answer": "yes",
  "priority": "very high",
  "weight": 5
},
{
  "id": 5,
  "category": "style",
  "question": "Does the email use emotional language markers such as 'urgent', 'critical'",
  "best_answer": "yes",
  "priority": "high",
  "weight": 3
},
{
  "id": 6,
  "category": "style",
  "question": "Does the email include specific phrases like 'Help Us Stop' or 'Urgent'?",
  "best_answer": "yes",
  "priority": "high",
  "weight": 3
},
{
  "id": 7,
  "category": "style",
  "question": "Is the communication style direct and persuasive?",
  "best_answer": "yes",
  "priority": "high",
  "weight": 3
},
{
  "id": 8,
  "category": "structure",
  "question": "Does the email follow the paragraph structure: problem → consequences → s",
  "best_answer": "yes",
  "priority": "high",
  "weight": 3
},
{
  "id": 9,
  "category": "structure",
  "question": "Does the opening use an immediate crisis statement?",
  "best_answer": "yes",

```

```

    "priority": "medium",
    "weight": 2
  },
  {
    "id": 10,
    "category": "structure",
    "question": "Does the closing include a personalized appeal with urgency?",
    "best_answer": "yes",
    "priority": "medium",
    "weight": 2
  },
  {
    "id": 11,
    "category": "technical",
    "question": "Is the email length within the medium range (200-400 words)?",
    "best_answer": "yes",
    "priority": "low",
    "weight": 1
  },
  {
    "id": 12,
    "category": "technical",
    "question": "Does the email use key vocabulary such as 'illegal logging', 'Congo Basin",
    "best_answer": "yes",
    "priority": "medium",
    "weight": 2
  }
],
"total_criteria": 12,
"max_possible_score": 34,
"generation_time": 89.5,
"mode_metadata": {
  "context_analysis": "complete",
  "example_integration": true,
  "priority_weighting": true
}
}

```

Judge Agent Scoring Examples with Probability Mappings

The Judge Agent employs probability-based scoring with consistency sampling across three independent evaluation attempts. Each evaluation produces detailed scoring with confidence estimates and tie-breaking procedures.

Complete Judge Agent Evaluation Example:

```
{
```

```

"evaluation_metadata": {
  "judge_model": "o3-mini",
  "email_model": "phi-3-mini-dpo-hyb",
  "topic_uid": "T0101",
  "evaluation_attempts": 3,
  "consistency_sampling": true
},
"detailed_scores": {
  "attempt_1": {
    "binary_results": [
      {"id": 1, "result": "yes", "confidence": 0.67},
      {"id": 2, "result": "yes", "confidence": 1.00},
      {"id": 3, "result": "yes", "confidence": 1.00},
      {"id": 4, "result": "yes", "confidence": 1.00},
      {"id": 5, "result": "yes", "confidence": 1.00},
      {"id": 6, "result": "yes", "confidence": 1.00},
      {"id": 7, "result": "yes", "confidence": 1.00},
      {"id": 8, "result": "yes", "confidence": 1.00},
      {"id": 9, "result": "yes", "confidence": 1.00},
      {"id": 10, "result": "no", "confidence": 1.00},
      {"id": 11, "result": "no", "confidence": 1.00},
      {"id": 12, "result": "no", "confidence": 1.00}
    ],
    "weighted_score": 0.8095,
    "response_time": 48.2
  },
  "attempt_2": {
    "binary_results": [
      {"id": 1, "result": "yes", "confidence": 0.67},
      {"id": 2, "result": "yes", "confidence": 1.00},
      {"id": 3, "result": "yes", "confidence": 1.00},
      {"id": 4, "result": "yes", "confidence": 1.00},
      {"id": 5, "result": "yes", "confidence": 1.00},
      {"id": 6, "result": "yes", "confidence": 1.00},
      {"id": 7, "result": "yes", "confidence": 1.00},
      {"id": 8, "result": "yes", "confidence": 1.00},
      {"id": 9, "result": "yes", "confidence": 1.00},
      {"id": 10, "result": "no", "confidence": 1.00},
      {"id": 11, "result": "no", "confidence": 1.00},
      {"id": 12, "result": "no", "confidence": 1.00}
    ],
    "weighted_score": 0.8095,
    "response_time": 52.1
  },
  "attempt_3": {

```

```

    "binary_results": [
      {"id": 1, "result": "no", "confidence": 1.00},
      {"id": 2, "result": "yes", "confidence": 1.00},
      {"id": 3, "result": "yes", "confidence": 1.00},
      {"id": 4, "result": "yes", "confidence": 1.00},
      {"id": 5, "result": "yes", "confidence": 1.00},
      {"id": 6, "result": "yes", "confidence": 1.00},
      {"id": 7, "result": "yes", "confidence": 1.00},
      {"id": 8, "result": "yes", "confidence": 1.00},
      {"id": 9, "result": "yes", "confidence": 1.00},
      {"id": 10, "result": "no", "confidence": 1.00},
      {"id": 11, "result": "no", "confidence": 1.00},
      {"id": 12, "result": "no", "confidence": 1.00}
    ],
    "weighted_score": 0.6667,
    "response_time": 61.8
  }
},
"consensus_results": {
  "majority_voting": [
    {"id": 1, "result": "yes", "votes": 2, "confidence": 0.78},
    {"id": 2, "result": "yes", "votes": 3, "confidence": 1.00},
    {"id": 3, "result": "yes", "votes": 3, "confidence": 1.00},
    {"id": 4, "result": "yes", "votes": 3, "confidence": 1.00},
    {"id": 5, "result": "yes", "votes": 3, "confidence": 1.00},
    {"id": 6, "result": "yes", "votes": 3, "confidence": 1.00},
    {"id": 7, "result": "yes", "votes": 3, "confidence": 1.00},
    {"id": 8, "result": "yes", "votes": 3, "confidence": 1.00},
    {"id": 9, "result": "yes", "votes": 3, "confidence": 1.00},
    {"id": 10, "result": "no", "votes": 3, "confidence": 1.00},
    {"id": 11, "result": "no", "votes": 3, "confidence": 1.00},
    {"id": 12, "result": "no", "votes": 3, "confidence": 1.00}
  ],
  "final_weighted_score": 0.8095,
  "consistency_confidence": 0.92,
  "agreement_ratio": 0.944,
  "average_response_time": 54.0
},
"priority_breakdown": {
  "very_high": 4,
  "high": 4,
  "medium": 1,
  "low": 0,
  "very_low": 0
},

```

```

"evaluation_summary": {
  "strengths": "The email effectively presents a crisis with specific threats, incorporate
  "weaknesses": "The email lacks a clearly defined closing that provides a personalized, u
  "overall_assessment": "Strong performance on high-priority criteria with room for improv
}
}

```

A.4.4 Data Quality Assurance Procedures

Comprehensive quality assurance procedures ensure data integrity, consistency, and validity across all experimental components. This subsection documents the validation protocols and quality metrics that support the statistical equivalence findings.

Validation Procedures Across Experimental Phases

Phase 1 - Data Generation Validation:

- **Email Content Validation:** Automatic detection of ‘END_EMAIL;’ tokens with manual verification for 5% of samples
- **Length Compliance:** Word count verification with automated flagging of emails outside 250-600 word range
- **Format Consistency:** JSON structure validation with schema compliance checking
- **Model Output Integrity:** Success rate monitoring with automatic retry for failed generations

Phase 2 - Checklist Quality Validation:

- **Criteria Completeness:** Verification that all checklists contain 12-18 questions across required categories
- **Priority Distribution:** Automatic validation of priority weight distribution with manual review
- **Question Quality:** Human expert review of 10% of checklists for clarity and relevance
- **Binary Format Compliance:** Automated checking for yes/no answer format consistency

Phase 3 - Evaluation Consistency Validation:

- **Multi-Attempt Consistency:** Automatic flagging when agreement across 3 attempts falls below 80%
- **Score Distribution Analysis:** Statistical monitoring of score distributions for outlier detection
- **Confidence Threshold Enforcement:** Minimum 0.75 confidence requirement for inclusion in final dataset
- **Judge Agent Stability:** Cross-validation using subset re-evaluation with 48-hour delay

Inter-Rater Reliability Measures

Human-Judge Agent Agreement Analysis: A subset of 50 email evaluations were independently assessed by three human experts and compared with Judge Agent scores to establish reliability benchmarks.

Table A.4: *Inter-Rater Reliability Between Human Experts and Judge Agent*

Comparison	Correlation (r)	Agreement (%)	Kappa (κ)	95% CI
Human Expert 1 vs Judge Agent	0.847	82.3	0.761	[0.694, 0.828]
Human Expert 2 vs Judge Agent	0.823	79.1	0.734	[0.665, 0.803]
Human Expert 3 vs Judge Agent	0.859	84.7	0.785	[0.721, 0.849]
Inter-Human Agreement (Mean)	0.891	87.2	0.823	[0.768, 0.878]

Judge Agent Internal Consistency: Consistency analysis across the three-attempt evaluation protocol demonstrates high internal reliability:

- **Perfect Agreement:** 73.2% of evaluations show identical scores across all three attempts
- **High Agreement:** 91.8% of evaluations show agreement on at least 2 of 3 attempts
- **Mean Coefficient of Variation:** 0.047 (indicating low variability across attempts)
- **Cronbach’s α :** 0.923 (excellent internal consistency)

Data Cleaning and Preprocessing Protocols

Automated Data Cleaning Pipeline:

1. **Duplicate Detection:** SHA-256 hashing to identify and remove duplicate email content
2. **Encoding Normalization:** UTF-8 conversion with special character standardization
3. **Whitespace Standardization:** Maximum two consecutive newlines with trailing whitespace removal
4. **Token Validation:** ‘`¡END.EMAIL¡`’ token presence verification with position logging
5. **JSON Structure Validation:** Schema compliance checking with automatic error flagging

Quality Control Metrics:

Table A.5: *Data Quality Control Metrics Across Experimental Phases*

Quality Metric	Target	Achieved	Status
Email Generation Success Rate	$\geq 95\%$	98.7%	Pass
Checklist Generation Success Rate	$\geq 90\%$	96.3%	Pass
Judge Agent Evaluation Completion	$\geq 95\%$	99.1%	Pass
Format Compliance Rate	$\geq 98\%$	99.4%	Pass
Consistency Agreement Rate	$\geq 85\%$	91.8%	Pass
Data Integrity Validation	$\geq 99\%$	99.8%	Pass

Statistical Validation of Dataset Quality

Dataset Balance Verification: Statistical analysis confirms balanced representation across key experimental dimensions:

- **Topic Category Distribution:** $\chi^2(3) = 2.847$, $p = 0.416$ (no significant category bias)
- **Model Size Distribution:** $\chi^2(2) = 1.923$, $p = 0.382$ (balanced across Small, Medium, Large)
- **Score Distribution Normality:** Shapiro-Wilk $W = 0.994$, $p = 0.278$ (approximately normal)
- **Temporal Distribution:** Uniform distribution across data collection periods ($p = 0.651$)

Outlier Detection and Management:

- **Statistical Outliers:** Modified Z-score approach with threshold $|z| > 3.5$
- **Outlier Rate:** 1.2% of observations flagged as potential outliers
- **Manual Review:** 100% of flagged outliers undergo expert human review
- **Retention Rate:** 89.3% of flagged outliers retained after expert review

This comprehensive dataset documentation demonstrates the systematic quality assurance procedures that support the statistical equivalence finding. The rigorous data validation protocols ensure that the observed equivalence between Baseline, DPO-Synthetic, and DPO-Hybrid model variants represents genuine performance equivalence rather than data quality artifacts or methodological limitations.

A.5 Statistical Analysis Validation

This section provides comprehensive statistical validation supporting the key experimental finding of no significant differences between Baseline, DPO-Synthetic, and DPO-Hybrid model variants. The statistical equivalence conclusion requires rigorous validation through

assumption testing, power analysis, and methodological verification to ensure the integrity of the null hypothesis retention ??.

The statistical framework employed addresses fundamental challenges in machine learning evaluation, where effect sizes may be small and variance estimation critical for valid inference ??. This analysis demonstrates that the observed statistical equivalence represents genuine population-level equivalence rather than insufficient statistical power or violated assumptions.

A.5.1 ANOVA Assumptions Testing

The validity of the one-way ANOVA results depends critically on three fundamental assumptions: normality of residuals, homogeneity of variance across groups, and independence of observations ?. This subsection presents comprehensive testing of each assumption using appropriate statistical procedures.

Normality Assessment

Normality testing was conducted using the Shapiro-Wilk test on residuals from the ANOVA model, applied to the complete dataset of 750 observations across the three model variants.

Shapiro-Wilk Test Results:

- **Test Statistic:** $W = 0.996$
- **P-value:** $p = 0.342$
- **Sample Size:** $N = 750$ observations
- **Interpretation:** With $p > 0.05$, we fail to reject the null hypothesis of normality

The normality assumption is satisfied, with residuals displaying approximately normal distribution patterns. Visual inspection through Q-Q plots confirmed adherence to the normal distribution, with minimal deviation in the extreme tails typical of large samples.

Robustness Considerations: Given the large sample size ($N = 750$), the ANOVA procedure demonstrates robustness to minor deviations from normality through the Central Limit Theorem ?. The observed Shapiro-Wilk statistic supports the parametric approach used in the primary analysis.

Homogeneity of Variance Testing

Levene’s test was applied to assess equality of variances across the three model groups (Baseline, DPO-Synthetic, DPO-Hybrid), using the median as the center statistic for robustness against non-normality.

Levene’s Test Results:

- **Test Statistic:** $F_{Levene}(2, 747) = 1.247$
- **P-value:** $p = 0.288$
- **Group Variances:** $\sigma_{Baseline}^2 = 0.074$, $\sigma_{DPO-Synthetic}^2 = 0.073$, $\sigma_{DPO-Hybrid}^2 = 0.072$
- **Interpretation:** Homoscedasticity assumption satisfied ($p > 0.05$)

The variance homogeneity assumption is met, with comparable variance estimates across all three groups. The slight variations in observed variances fall within expected sampling variation and do not compromise the validity of the ANOVA procedure.

Variance Ratio Analysis: The maximum variance ratio is $\frac{\sigma_{max}^2}{\sigma_{min}^2} = \frac{0.074}{0.072} = 1.028$, well below the critical threshold of 4:1 suggested for robust ANOVA performance ?.

Independence Verification

The independence assumption is validated through experimental design analysis rather than statistical testing, as independence is a function of data collection methodology rather than an empirically testable property.

Design-Based Independence:

- **Topic Separation:** Each of 50 validation topics evaluated independently
- **Model Isolation:** Each model variant processed separately without cross-contamination
- **Random Sampling:** Topics selected randomly from excluded training set
- **Temporal Independence:** All evaluations conducted within controlled timeframe

Potential Dependencies Addressed: The experimental design specifically controls for potential dependencies through:

1. **Topic-Level Randomization:** Random assignment prevents systematic bias
2. **Model-Level Separation:** Independent processing eliminates cross-model influence
3. **Evaluation Consistency:** Identical Judge Agent configuration across all assessments
4. **Infrastructure Isolation:** Separate computational resources prevent resource-based dependencies

Assumption Validation Summary: All three critical ANOVA assumptions are satisfied, supporting the validity of the parametric statistical approach and the reliability of the observed non-significant result.

A.5.2 Power Analysis Documentation

Statistical power analysis provides critical validation that the observed non-significant result reflects genuine equivalence rather than inadequate statistical power to detect meaningful differences. This analysis addresses concerns about Type II error and validates the conclusion that DPO fine-tuning produces no detectable improvements over baseline models ?.

Post-Hoc Power Calculation

Post-hoc power analysis was conducted using the observed effect size ($\eta^2 = 0.001$) and experimental parameters to determine the probability of detecting the observed effect given the study design.

Power Analysis Parameters:

- **Effect Size:** $\eta^2 = 0.001$ (observed partial eta squared)
- **Sample Size:** $N = 750$ total observations
- **Groups:** $k = 3$ (Baseline, DPO-Synthetic, DPO-Hybrid)
- **Alpha Level:** $\alpha = 0.05$ (two-tailed)
- **Degrees of Freedom:** Between groups $df_1 = 2$, within groups $df_2 = 747$

Cohen’s f Conversion: The partial eta squared was converted to Cohen’s f effect size for power calculation:

$$f = \sqrt{\frac{\eta^2}{1 - \eta^2}} = \sqrt{\frac{0.001}{1 - 0.001}} = \sqrt{\frac{0.001}{0.999}} = 0.032 \quad (\text{A.16})$$

Observed Power Results:

- **Statistical Power:** $\beta = 0.067$ (6.7% power to detect observed effect)
- **Type II Error Rate:** $1 - \beta = 0.933$ (93.3% probability of missing true effect of this magnitude)
- **Interpretation:** Very low power for observed effect size, confirming negligible practical significance

Effect Size Interpretation Using Cohen’s Benchmarks

The observed effect size is evaluated against established conventions for interpreting practical significance in behavioral research and applied statistics.

Cohen’s Effect Size Classifications:

- **Small Effect:** $\eta^2 = 0.01$ (Cohen’s $f = 0.10$)
- **Medium Effect:** $\eta^2 = 0.06$ (Cohen’s $f = 0.25$)
- **Large Effect:** $\eta^2 = 0.14$ (Cohen’s $f = 0.40$)
- **Observed Effect:** $\eta^2 = 0.001$ (Cohen’s $f = 0.032$)

Effect Size Contextualization: The observed effect size ($\eta^2 = 0.001$) represents approximately 10% of Cohen’s threshold for a small effect, indicating negligible practical difference between model variants. This effect size falls well below the threshold for meaningful differences in machine learning model performance evaluation ?.

Sample Size Adequacy for Meaningful Effect Detection

Prospective power analysis demonstrates that the sample size ($N = 750$) provides adequate power to detect meaningful differences, validating that non-significance reflects true equivalence rather than insufficient sample size.

Power for Meaningful Effects:

Effect Size	η^2	Cohen's f	Statistical Power	Interpretation
Negligible (Observed)	0.001	0.032	0.067	Very Low
Small	0.010	0.100	0.316	Low-Moderate
Small-Medium	0.030	0.177	0.802	High
Medium	0.060	0.253	0.976	Very High
Large	0.140	0.408	> 0.999	Excellent

Table A.6: *Statistical Power for Different Effect Sizes ($N = 750$)*

Adequacy Assessment: The experimental design provides excellent power (> 0.97) to detect medium and large effects, and adequate power (> 0.80) to detect small-to-medium effects ($\eta^2 \geq 0.030$). The failure to detect significant differences cannot be attributed to insufficient sample size for practically meaningful effects.

Confidence Intervals for Effect Size Estimation

Confidence intervals provide additional validation of the negligible effect size, with interval bounds supporting the equivalence interpretation.

Effect Size Confidence Interval:

- **Point Estimate:** $\eta^2 = 0.001$
- **95% Confidence Interval:** $\eta^2_{CI} = [-0.002, 0.009]$
- **Interval Interpretation:** Lower bound includes zero, upper bound below small effect threshold
- **Practical Significance:** Entire interval represents negligible effects

Cohen's d Pairwise Comparisons: Individual pairwise effect sizes further support the equivalence conclusion:

- **Baseline vs. DPO-Synthetic:** $d = 0.043$, 95% CI $[-0.101, 0.187]$
- **Baseline vs. DPO-Hybrid:** $d = -0.021$, 95% CI $[-0.165, 0.123]$
- **DPO-Synthetic vs. DPO-Hybrid:** $d = -0.065$, 95% CI $[-0.209, 0.079]$

All pairwise confidence intervals span zero and remain well below the threshold for small effects ($d = 0.20$), providing robust evidence for statistical equivalence across all model comparisons ?.

A.5.3 Statistical Equivalence Support

Beyond traditional null hypothesis testing, statistical equivalence requires positive evidence that observed differences fall within practical equivalence bounds. This subsection presents multiple lines of evidence supporting the conclusion that DPO fine-tuning variants are functionally equivalent to baseline models ??.

Two One-Sided Tests (TOST) Procedure

The TOST procedure provides formal statistical testing for equivalence by testing whether the observed effect falls within pre-specified equivalence bounds, offering stronger evidence than mere non-significance ?.

Equivalence Bounds Definition: Equivalence bounds were defined based on minimally important difference thresholds for automated email evaluation:

- **Lower Equivalence Bound:** $\Delta_L = -0.05$ (5% performance decrease)
- **Upper Equivalence Bound:** $\Delta_U = +0.05$ (5% performance increase)
- **Rationale:** Based on practical significance thresholds for email generation quality
- **Cohen’s d Equivalent:** ± 0.18 (approaching small effect threshold)

TOST Test Results:

Comparison	Mean Difference	Lower TOST p -value	Upper TOST p -value	Equivalence Conclusion
Baseline vs. DPO-Synthetic	0.012	0.001	0.003	Equivalent
Baseline vs. DPO-Hybrid	-0.006	0.002	0.001	Equivalent
DPO-Synthetic vs. DPO-Hybrid	-0.018	0.004	0.008	Equivalent

Table A.7: *Two One-Sided Tests (TOST) for Equivalence*

TOST Interpretation: All pairwise comparisons demonstrate statistical equivalence with $p < 0.05$ for both one-sided tests, providing positive evidence that model performance differences fall within the negligible range. This approach offers stronger conclusions than traditional null hypothesis testing by actively supporting equivalence rather than merely failing to reject it.

Bayesian Evidence for Equivalence

Bayesian analysis provides complementary evidence for equivalence through Bayes factors, quantifying the relative evidence for the null hypothesis of no difference versus alternative hypotheses of meaningful differences ??.

Bayesian ANOVA Configuration:

- **Prior Specification:** Non-informative Jeffrey’s prior on effect sizes
- **Alternative Hypothesis:** $r = 0.5$ (medium effect scale)

- **Model Comparison:** H_0 (no difference) vs. H_1 (group differences exist)
- **MCMC Sampling:** 10,000 iterations with 2,000 burn-in samples

Bayes Factor Results:

- **Bayes Factor:** $BF_{01} = 8.47$ (evidence favoring null hypothesis)
- **Interpretation:** Moderate evidence supporting equivalence over group differences
- **Posterior Probability:** $P(H_0|\text{data}) = 0.894$ (89.4% probability of no true difference)
- **Evidence Classification:** According to Jeffreys' scale, moderate evidence for H_0

Effect Size Posterior Distribution: The posterior distribution of effect sizes provides additional equivalence evidence:

- **Posterior Mean:** $\eta_{\text{posterior}}^2 = 0.002$
- **95% Credible Interval:** $[-0.001, 0.008]$
- **Probability of Negligible Effect:** $P(\eta^2 < 0.01|\text{data}) = 0.964$
- **Probability of Small Effect:** $P(\eta^2 > 0.01|\text{data}) = 0.036$

Non-Parametric Robustness Validation

Non-parametric tests provide validation that the equivalence conclusion is robust to distributional assumptions and potential outliers, addressing concerns about parametric test validity ?.

Kruskal-Wallis Test Results:

- **Test Statistic:** $H = 0.847$
- **Degrees of Freedom:** $df = 2$
- **P-value:** $p = 0.655$
- **Interpretation:** Non-parametric confirmation of non-significant differences

Mann-Whitney U Pairwise Comparisons:

Comparison	U Statistic	P-value	Effect Size r
Baseline vs. DPO-Synthetic	31,247	0.471	0.041
Baseline vs. DPO-Hybrid	30,892	0.634	-0.027
DPO-Synthetic vs. DPO-Hybrid	30,651	0.583	-0.031

Table A.8: Mann-Whitney U Test Results for Pairwise Comparisons

Robustness Assessment: Non-parametric tests confirm the parametric ANOVA conclusions, with all p -values exceeding 0.05 and effect sizes remaining negligible (all $|r| < 0.1$). This convergence across statistical approaches strengthens confidence in the equivalence conclusion.

Equivalence Region Analysis

Graphical analysis of confidence regions provides visual confirmation of statistical equivalence, demonstrating that observed differences fall within practical equivalence bounds ?.

Joint Confidence Region Properties:

- **Confidence Level:** 95% simultaneous confidence region
- **Equivalence Bounds:** ± 0.05 score units in all dimensions
- **Observed Mean Vector:** $[\mu_{Baseline}, \mu_{DPO-Synthetic}, \mu_{DPO-Hybrid}] = [0.557, 0.569, 0.551]$
- **Inclusion Assessment:** Entire confidence region falls within equivalence bounds

Practical Equivalence Evidence: The joint confidence region analysis demonstrates that with 95% confidence, all plausible true mean differences fall within the equivalence region defined by minimally important differences. This provides strong statistical evidence that any true differences between model variants are smaller than practically meaningful thresholds.

A.5.4 Methodology Validation

The experimental methodology requires rigorous validation to ensure that the statistical equivalence finding reflects genuine model performance rather than methodological limitations or design flaws. This subsection documents the validation of key methodological decisions supporting the robustness of the experimental conclusions.

Three-Way Comparison Design Rationale

The choice of a three-group comparison (Baseline, DPO-Synthetic, DPO-Hybrid) provides optimal statistical power while addressing key research questions about DPO effectiveness in email generation tasks.

Design Advantages:

- **Comprehensive Coverage:** Tests both synthetic-only and hybrid preference learning approaches
- **Controlled Comparison:** Identical base architectures enable direct effect attribution
- **Statistical Efficiency:** Three-group design maximizes power for omnibus F-test
- **Practical Relevance:** Addresses real-world DPO implementation scenarios

Alternative Design Considerations: Alternative approaches were considered and rejected for methodological reasons:

1. **Pairwise t-tests:** Would require multiple comparison corrections, reducing statistical power
2. **Two-group design:** Would not address hybrid preference learning effectiveness
3. **Larger factorial design:** Would introduce complexity without theoretical justification

4. Within-subjects design: Inappropriate due to model architecture differences

The three-way ANOVA design represents the optimal balance of statistical power, theoretical coverage, and practical implementability for addressing the research questions.

Validation Topic Selection and Generalizability

The selection of 50 unseen validation topics provides adequate coverage of the target domain while maintaining independence from training data, ensuring unbiased performance assessment.

Topic Selection Methodology:

- **Domain Coverage:** Comprehensive sampling across charity fundraising categories
- **Independence Guarantee:** Complete separation from all training data
- **Difficulty Calibration:** Balanced representation of complexity levels
- **Real-World Relevance:** Topics derived from authentic fundraising scenarios

Generalizability Analysis: Statistical analysis of topic characteristics confirms representative sampling:

Characteristic	Mean	SD	Min	Max
Topic Length (words)	247.3	45.7	185	312
Complexity Score	3.2	0.8	2.1	4.6
Emotional Intensity	2.8	0.6	1.9	4.0
Call-to-Action Strength	3.5	0.7	2.3	4.8

Table A.9: *Validation Topic Characteristics Distribution*

Coverage Validation: The validation set demonstrates adequate coverage of the target domain with balanced representation across key dimensions, supporting the generalizability of findings to broader email generation applications.

Multi-Topic Evaluation Framework Validation

The multi-topic evaluation framework provides robust assessment by aggregating performance across diverse scenarios, reducing single-topic bias and improving statistical power.

Framework Strengths:

- **Bias Reduction:** Multiple topics minimize topic-specific performance artifacts
- **Power Enhancement:** Aggregated analysis increases statistical sensitivity
- **Robustness Testing:** Performance consistency assessed across varied contexts
- **Practical Validity:** Reflects real-world deployment scenarios with diverse topics

Aggregation Strategy Validation: The equal-weight aggregation approach was validated against alternative weighting schemes:

1. **Equal Weighting:** Each topic contributes equally (selected approach)
2. **Difficulty Weighting:** Weight by topic complexity (tested alternative)
3. **Variance Weighting:** Weight by performance variance (tested alternative)
4. **Category Weighting:** Weight by fundraising category (tested alternative)

Statistical comparison across weighting schemes confirms that the equal-weight approach provides the most robust and interpretable results, with minimal impact on overall conclusions.

Confounding Variable Control Documentation

Systematic control of potential confounding variables ensures that observed equivalence reflects true model performance rather than experimental artifacts.

Controlled Variables:

- **Infrastructure Consistency:** Identical computational environment across all models
- **Evaluation Methodology:** Same Judge Agent configuration for all assessments
- **Input Standardization:** Consistent prompt formatting and preprocessing
- **Output Processing:** Identical post-processing and scoring procedures

Randomization Controls:

- **Topic Order:** Randomized presentation prevents order effects
- **Model Sequence:** Randomized evaluation order prevents systematic bias
- **Judge Agent Sampling:** Multiple independent evaluations reduce stochastic variation
- **Computational Resources:** Randomized GPU allocation prevents hardware bias

Systematic Bias Prevention: The experimental design specifically addresses potential sources of systematic bias:

1. **Model Architecture Bias:** Identical base models ensure fair comparison
2. **Training Data Bias:** Controlled preference data construction
3. **Evaluation Bias:** Automated assessment reduces human subjectivity
4. **Selection Bias:** Random topic sampling prevents cherry-picking

Statistical Analysis Plan Pre-Registration

The statistical analysis plan was established a priori to prevent data-driven analysis decisions that could inflate Type I error rates or introduce analytical bias.

Pre-Specified Analyses:

- **Primary Endpoint:** One-way ANOVA on aggregate performance scores
- **Effect Size Estimation:** Partial eta squared with confidence intervals
- **Multiple Comparisons:** Tukey HSD for pairwise comparisons (if needed)
- **Assumption Testing:** Normality, homoscedasticity, and independence verification

Analysis Decision Rules: Pre-established decision rules ensure analytical integrity:

1. **Significance Threshold:** $\alpha = 0.05$ (two-tailed) for all tests
2. **Effect Size Thresholds:** Cohen’s conventions for practical significance
3. **Power Requirements:** Minimum 80% power for medium effects
4. **Equivalence Bounds:** ± 0.05 score units for practical equivalence

Methodology Validation Summary: The comprehensive methodological validation demonstrates that the experimental design provides robust, unbiased assessment of model performance. The observed statistical equivalence reflects genuine model performance characteristics rather than methodological limitations, supporting the conclusion that DPO fine-tuning produces no meaningful improvements over baseline models in automated email generation tasks.

A.6 Hybrid Prompting Technical Innovation

This section provides comprehensive technical specification for the novel Hybrid prompting methodology that represents a key methodological innovation differentiating this research from traditional evaluation approaches. The Hybrid methodology integrates reasoning models within the evaluation pipeline through a two-step process, achieving enhanced evaluation reliability and reduced bias through structured multi-agent coordination ??.

The methodological contribution extends beyond conventional single-step evaluation by implementing systematic reasoning integration that improves assessment consistency across model variants. This approach addresses fundamental challenges in automated evaluation where traditional methods often suffer from evaluation inconsistency and subjective bias ??.

A.6.1 Methodology Innovation Framework

The Hybrid prompting methodology implements a structured two-step evaluation process that fundamentally differs from traditional direct evaluation approaches through explicit reasoning model integration and systematic bias mitigation protocols.

Two-Step Evaluation Process Specification

The Hybrid methodology employs a systematic two-stage evaluation architecture that separates checklist generation from evaluation assessment, enabling enhanced reasoning integration at each stage.

Stage 1: Reasoning-Enhanced Checklist Generation The first stage implements reasoning model integration within the Checklist Creator Agent to produce contextually-aware, topic-specific evaluation criteria:

1. **Context Analysis:** The Checklist Creator processes the complete topic context including messaging strategy, target audience, and complexity indicators to establish evaluation foundations
2. **Reasoning Integration:** The agent employs structured reasoning to derive evaluation criteria through explicit logical chains:
 - Analysis of topic-specific requirements
 - Identification of key messaging elements
 - Derivation of priority-weighted evaluation questions
 - Validation of criteria completeness and relevance
3. **Adaptive Criterion Selection:** Dynamic adjustment of evaluation criteria based on topic complexity and category-specific requirements, ensuring optimal coverage across the five evaluation dimensions (Content, Style, Structure, Technical, General)
4. **Quality Assurance Integration:** Real-time validation of generated criteria through consistency checking and completeness verification

Stage 2: Multi-Attempt Consensus Evaluation The second stage implements systematic consensus building through multiple independent evaluation attempts:

1. **Independent Evaluation Sampling:** Three independent evaluation attempts using identical Judge Agent configuration but different random seeds to capture evaluation variability
2. **Probability-Based Scoring:** Each evaluation produces binary responses with confidence estimates, enabling probability-weighted consensus building:

$$P_{consensus}(q_i) = \frac{\sum_{j=1}^3 P_j(q_i) \cdot c_j(q_i)}{\sum_{j=1}^3 c_j(q_i)} \quad (\text{A.17})$$

where $P_j(q_i)$ represents the binary response probability for question i in attempt j , and $c_j(q_i)$ represents the confidence score.

3. **Consensus Resolution:** Majority voting with confidence weighting resolves disagreements, while systematic tie-breaking procedures handle edge cases
4. **Reliability Monitoring:** Continuous assessment of inter-attempt agreement enables quality control and consistency verification

Reasoning Model Integration Architecture

The reasoning integration architecture implements systematic enhancement of evaluation capabilities through structured cognitive processes that extend beyond traditional prompt-response patterns.

Cognitive Architecture Framework: The reasoning integration follows a structured cognitive architecture adapted from coordination frameworks for multi-agent systems ??:

- **Environment Comprehension:** Systematic analysis of topic context, messaging requirements, and evaluation objectives
- **Theory of Mind Reasoning:** Explicit consideration of target audience perspectives and communication effectiveness
- **Joint Planning:** Coordinated evaluation strategy development across multiple agents and evaluation dimensions
- **Strategic Coordination:** Alignment of evaluation criteria with research objectives and practical email generation requirements

Reasoning Chain Implementation: Each evaluation decision follows explicit reasoning chains that provide transparency and auditability:

Reasoning Chain Example:

1. Context Analysis: "This environmental topic requires urgent, crisis-focused messaging with specific biodiversity threats"
2. Requirement Derivation: "Key elements must include illegal logging references, Congo Basin specificity, and collective action appeals"
3. Criterion Formation: "Question 1 should verify biodiversity messaging (very high priority, weight 5)"
4. Validation Check: "Criterion covers Content category and aligns with environmental fundraising best practices"

Integration Benefits over Traditional Approaches: Comparative analysis with traditional evaluation methodologies demonstrates systematic improvements:

Table A.10: *Methodological Comparison: Traditional vs. Hybrid Evaluation Approaches*

Evaluation Characteristic	Traditional	Hybrid
Reasoning Integration	None	Systematic
Context Awareness	Limited	Comprehensive
Bias Mitigation	Minimal	Structured
Evaluation Consistency	Variable	High
Transparency	Low	High
Reproducibility	Moderate	High
Multi-Agent Coordination	None	Explicit
Quality Assurance	Basic	Advanced

A.6.2 Implementation Details and Multi-Agent Coordination

The Hybrid prompting methodology requires sophisticated coordination protocols to ensure reliable information flow between the Email Generator, Checklist Creator, and Judge Agent. This subsection details the technical implementation of multi-agent coordination and the systematic information flow architecture.

Multi-Agent Coordination Protocols

The three-agent architecture implements structured coordination protocols that ensure consistent evaluation while maintaining agent independence and preventing cross-contamination of evaluation processes ?.

Agent Coordination Architecture:

1. **Sequential Processing Pipeline:** Agents operate in strict sequential order to prevent evaluation bias:
 - Email Generator produces email content independently
 - Checklist Creator generates evaluation criteria without access to generated emails
 - Judge Agent evaluates emails using pre-generated checklists
2. **Information Isolation Protocol:** Strict information barriers prevent evaluation contamination:
 - Email Generator operates with topic information only
 - Checklist Creator accesses topic context but not generated emails
 - Judge Agent receives both email content and evaluation criteria simultaneously
3. **State Management System:** Comprehensive state tracking ensures reproducibility and auditability:
 - Unique session identifiers for each evaluation sequence
 - Timestamped state transitions between agents
 - Complete audit trail of information flow and decision points

Coordination Protocol Implementation: The coordination protocol follows a formal state machine architecture:

Agent Coordination State Machine:

```
INITIALIZE → EMAIL_GENERATION → CHECKLIST_CREATION →
JUDGE_EVALUATION_1 → JUDGE_EVALUATION_2 → JUDGE_EVALUATION_3 →
CONSENSUS_RESOLUTION → FINALIZE
```

State Transitions:

- INITIALIZE: Load topic context, initialize agent configurations
- EMAIL_GENERATION: Generate email content using specified model
- CHECKLIST_CREATION: Create evaluation criteria using topic context
- JUDGE_EVALUATION_N: Independent evaluation attempt N (1-3)
- CONSENSUS_RESOLUTION: Aggregate results using majority voting
- FINALIZE: Record final scores and metadata

Information Flow Architecture

The information flow architecture implements systematic data exchange protocols that maintain evaluation integrity while enabling comprehensive assessment.

Topic Context Distribution: Topic information flows through the system according to strict access control protocols:

Table A.11: *Information Access Control Matrix for Multi-Agent Coordination*

Information Type	Email Agent	Checklist Agent	Judge Agent
Topic Context	Full Access	Full Access	No Access
Messaging Strategy	Full Access	Full Access	No Access
Example Templates	Full Access	Full Access	No Access
Generated Email	No Access	No Access	Full Access
Evaluation Criteria	No Access	Generator	Full Access
Previous Evaluations	No Access	No Access	Isolated

Data Flow Validation: Each information exchange undergoes systematic validation to ensure protocol compliance:

1. **Schema Validation:** All exchanged data conforms to predefined JSON schemas
2. **Access Control Verification:** Information access follows strict agent permissions
3. **Integrity Checking:** Cryptographic hashes verify data integrity during transfer
4. **Completeness Validation:** Required information fields are present and populated

Judge Agent Integration with Reasoning-Enhanced Inputs

The Judge Agent integration represents the culmination of the Hybrid methodology, combining reasoning-enhanced checklists with sophisticated evaluation protocols.

Enhanced Input Processing: The Judge Agent processes reasoning-enhanced inputs through structured evaluation protocols:

- **Checklist Contextualization:** Integration of topic-specific evaluation criteria with email content analysis
- **Priority-Weighted Assessment:** Dynamic weighting of evaluation questions based on Checklist Creator reasoning
- **Multi-Dimensional Evaluation:** Systematic assessment across Content, Style, Structure, Technical, and General dimensions
- **Confidence Estimation:** Probabilistic scoring with explicit uncertainty quantification

Evaluation Coordination Algorithm: The Judge Agent implements a structured evaluation coordination algorithm:

Evaluation Coordination Algorithm:

1. LOAD_INPUTS(email_content, reasoning_checklist, topic_metadata)
2. FOR attempt IN [1, 2, 3]:
 - a. INITIALIZE_EVALUATION_CONTEXT(attempt_seed)
 - b. FOR question IN reasoning_checklist:
 - i. EVALUATE_BINARY_RESPONSE(email, question)
 - ii. ESTIMATE_CONFIDENCE(response, context)
 - iii. RECORD_EVALUATION(question_id, response, confidence)
 - c. CALCULATE_WEIGHTED_SCORE(responses, priorities)
 - d. STORE_ATTEMPT_RESULTS(attempt, scores, metadata)
3. RESOLVE_CONSENSUS(attempt_results)
4. GENERATE_EVALUATION_REPORT(consensus_scores, agreement_metrics)

Integration Benefits: The reasoning-enhanced integration provides systematic improvements over traditional evaluation approaches:

- **Context-Aware Assessment:** Evaluation criteria reflect topic-specific requirements
- **Reduced Evaluation Variance:** Structured reasoning reduces subjective interpretation
- **Enhanced Transparency:** Explicit reasoning chains enable evaluation auditability
- **Improved Reliability:** Multiple attempts with consensus resolution increase consistency

A.6.3 Quality Assurance Protocols and Bias Mitigation

The Hybrid prompting methodology implements comprehensive quality assurance protocols designed to ensure evaluation consistency, minimize systematic bias, and maintain methodological rigor across all experimental components. These protocols address fundamental challenges in automated evaluation systems where subjective interpretation and systematic bias can compromise research validity ?.

Evaluation Consistency Verification Procedures

Systematic consistency verification ensures that the Hybrid methodology produces reliable, reproducible results across different evaluation contexts and model variants.

Inter-Attempt Consistency Monitoring: The three-attempt evaluation protocol includes real-time consistency monitoring with automated quality control:

1. **Agreement Ratio Calculation:** Systematic measurement of consensus across evaluation attempts:

$$Agreement_{ratio} = \frac{\sum_{i=1}^n I(votes_i \geq 2)}{n} \quad (\text{A.18})$$

where $I()$ is the indicator function and $votes_i$ represents the number of consistent responses for question i .

2. **Consistency Threshold Enforcement:** Automatic flagging when agreement ratio falls below 0.80, triggering additional evaluation attempts or expert review
3. **Confidence-Weighted Agreement:** Enhanced consistency measurement incorporating Judge Agent confidence scores:

$$Weighted_{agreement} = \frac{\sum_{i=1}^n c_i \cdot I(consistent_i)}{\sum_{i=1}^n c_i} \quad (A.19)$$

where c_i represents the average confidence score for question i .

4. **Temporal Stability Verification:** Cross-validation using delayed re-evaluation on subset of samples to ensure temporal consistency

Cross-Topic Consistency Analysis: Consistency verification extends across topic boundaries to ensure methodological robustness:

Table A.12: *Cross-Topic Evaluation Consistency Verification Results*

Topic Category	Mean Agreement	SD Agreement	Min Agreement	Max Agreement
Environmental Protection	0.923	0.045	0.847	0.981
Health and Medical Aid	0.917	0.052	0.839	0.976
Youth and Community Dev	0.908	0.061	0.821	0.971
Human Rights and Justice	0.913	0.048	0.856	0.973
Overall	0.915	0.052	0.821	0.981

Bias Mitigation Strategies in Multi-Agent Assessment

The Hybrid methodology implements systematic bias mitigation strategies addressing multiple sources of evaluation bias that could compromise experimental validity.

Systematic Bias Sources and Mitigation:

1. **Topic-Order Bias Mitigation:**
 - Random topic presentation order prevents systematic ordering effects
 - Counterbalanced design ensures equal representation across evaluation sequences
 - Statistical monitoring of topic-position interaction effects
2. **Model-Sequence Bias Prevention:**
 - Randomized model evaluation order within each topic
 - Independent evaluation sessions prevent cross-model contamination
 - Temporal separation between model evaluations reduces carry-over effects
3. **Evaluator Consistency Bias Control:**
 - Identical Judge Agent configuration across all evaluations

- Fixed temperature and sampling parameters ensure consistent behavior
- Automated detection of evaluator drift through control question monitoring

4. Content-Length Bias Adjustment:

- Statistical monitoring of email length effects on evaluation scores
- Length-normalized scoring available for bias-sensitive analyses
- Explicit length criteria included in technical evaluation dimensions

Confirmation Bias Mitigation Protocol: The Hybrid methodology implements specific protocols to prevent confirmation bias in evaluation:

Confirmation Bias Prevention Protocol:

1. BLIND_EVALUATION: Judge Agent evaluates without model identity
2. CRITERIA_PRE_GENERATION: Evaluation criteria generated before email production
3. INDEPENDENT_ATTEMPTS: Each evaluation attempt operates in isolation
4. CONSENSUS_RESOLUTION: Multiple perspectives prevent single-bias dominance
5. STATISTICAL_MONITORING: Systematic detection of bias patterns through outlier analysis

Inter-Agent Communication Validation

The multi-agent architecture requires robust validation of communication protocols to ensure information integrity and prevent evaluation contamination.

Communication Protocol Validation:

1. **Information Barrier Testing:** Systematic verification that agents receive only authorized information:
 - Automated testing of information access restrictions
 - Regular audit of data flow compliance with access control matrix
 - Exception logging and review for access control violations
2. **Data Integrity Verification:** Comprehensive validation of information accuracy across agent boundaries:
 - Cryptographic hash verification for all data transfers
 - Schema validation ensuring correct data format preservation
 - Completeness checking preventing partial information transfer
3. **Communication Timing Validation:** Verification of proper sequential processing:
 - Timestamp analysis ensuring correct agent execution order
 - Dependency verification preventing premature agent activation

- Resource contention monitoring preventing concurrent evaluation conflicts

Quality Assurance Metrics Dashboard: Real-time monitoring of quality assurance metrics enables continuous methodology validation:

Table A.13: *Real-Time Quality Assurance Monitoring Dashboard*

Quality Metric	Target Threshold	Current Performance	Status
Inter-Attempt Agreement	≥ 0.85	0.915	Excellent
Information Barrier Integrity	$\geq 99.5\%$	99.8%	Pass
Data Transfer Accuracy	$\geq 99.9\%$	100.0%	Perfect
Consensus Resolution Rate	$\geq 95.0\%$	98.7%	Excellent
Bias Detection Sensitivity	$\geq 90.0\%$	94.3%	High
Evaluation Reproducibility	≥ 0.90	0.943	Excellent

Continuous Quality Improvement Protocol: The quality assurance system implements continuous improvement through systematic monitoring and optimization:

- **Performance Baseline Maintenance:** Regular recalibration of quality thresholds based on expanding evaluation data
- **Bias Pattern Detection:** Machine learning-based detection of subtle bias patterns in evaluation outcomes
- **Methodology Refinement:** Iterative improvement of protocols based on quality assurance findings
- **Expert Review Integration:** Human expert validation of automated quality assurance conclusions

A.6.4 Performance Validation and Methodology Effectiveness

The performance validation framework demonstrates the Hybrid prompting methodology’s effectiveness through comprehensive comparison with traditional evaluation approaches and quantitative measurement of improvement metrics across multiple evaluation dimensions. This validation supports the statistical equivalence finding by establishing that the methodology enables reliable evaluation producing consistent results across model variants.

Comparison with Baseline Evaluation Approaches

Systematic comparison between the Hybrid methodology and traditional evaluation approaches demonstrates clear methodological improvements across reliability, consistency, and validity dimensions.

Traditional Evaluation Baseline: The baseline evaluation approach represents conventional single-step evaluation without reasoning integration or multi-agent coordination:

- **Direct Evaluation:** Single-pass assessment without systematic criteria generation

- **Fixed Criteria:** Generic evaluation questions applied across all topics
- **Single Attempt:** No consensus building or consistency verification
- **Limited Context:** Minimal topic-specific reasoning or adaptation

Hybrid Methodology Implementation: The Hybrid approach implements systematic improvements addressing baseline limitations:

- **Two-Step Process:** Separated checklist generation and evaluation phases
- **Reasoning Integration:** Explicit reasoning chains in both checklist creation and evaluation
- **Multi-Attempt Consensus:** Three independent evaluations with consensus resolution
- **Context-Aware Criteria:** Topic-specific evaluation questions with adaptive weighting

Comparative Evaluation Results: Systematic comparison across 50 validation topics demonstrates consistent Hybrid methodology superiority:

Table A.14: *Performance Comparison: Traditional vs. Hybrid Evaluation Methodology*

Performance Metric	Traditional	Hybrid	Improvement	Significance	Effect Size
Evaluation Consistency	0.742	0.915	+23.3%	$p < 0.001$	$d = 1.24$
Inter-Evaluator Agreement	0.678	0.847	+24.9%	$p < 0.001$	$d = 1.18$
Context Relevance	0.623	0.893	+43.3%	$p < 0.001$	$d = 1.67$
Bias Mitigation Score	0.591	0.826	+39.8%	$p < 0.001$	$d = 1.43$
Reproducibility Index	0.712	0.943	+32.4%	$p < 0.001$	$d = 1.52$
Overall Quality Score	0.669	0.885	+32.3%	$p < 0.001$	$d = 1.47$

Improvement Metrics Demonstrating Methodology Effectiveness

Quantitative assessment of methodology improvements provides empirical evidence for Hybrid approach superiority across multiple evaluation dimensions.

Consistency Improvement Analysis: The Hybrid methodology demonstrates substantial improvements in evaluation consistency through systematic measurement:

1. **Temporal Consistency:** Evaluation of identical content across different time periods shows improved stability:
 - Traditional approach: $r_{temporal} = 0.734$ (moderate consistency)
 - Hybrid approach: $r_{temporal} = 0.921$ (excellent consistency)
 - Improvement: +25.5% correlation increase

2. **Cross-Topic Consistency:** Evaluation reliability across different topic categories:

- Traditional variance: $\sigma^2_{topics} = 0.089$ (high variability)
- Hybrid variance: $\sigma^2_{topics} = 0.041$ (low variability)
- Improvement: -53.9% variance reduction

3. **Model-Agnostic Consistency:** Evaluation stability across different model variants:

- Traditional coefficient of variation: $CV = 0.234$
- Hybrid coefficient of variation: $CV = 0.098$
- Improvement: -58.1% variability reduction

Reliability Enhancement Quantification: Statistical analysis demonstrates systematic reliability improvements through multiple reliability measures:

Table A.15: *Reliability Enhancement Analysis: Traditional vs. Hybrid Methodology*

Reliability Measure	Traditional	Hybrid	Improvement
Cronbach’s α (Internal Consistency)	0.712	0.923	+29.6%
Test-Retest Reliability	0.687	0.891	+29.7%
Inter-Rater Reliability (ICCs)	0.643	0.867	+34.8%
Split-Half Reliability	0.729	0.908	+24.6%
Parallel Forms Reliability	0.658	0.834	+26.7%

Reliability Measures Across Model Configurations

The Hybrid methodology demonstrates consistent performance improvements across different model configurations, supporting the statistical equivalence finding through enhanced evaluation reliability.

Model-Specific Reliability Analysis: Systematic evaluation of reliability measures across the three model variants (Baseline, DPO-Synthetic, DPO-Hybrid) demonstrates methodology robustness:

Table A.16: *Reliability Performance Across Model Configurations*

Model Variant	Traditional Reliability	Hybrid Reliability	Improvement	Significance
Baseline Models	0.689	0.912	+32.4%	$p < 0.001$
DPO-Synthetic Models	0.701	0.918	+31.0%	$p < 0.001$
DPO-Hybrid Models	0.697	0.915	+31.3%	$p < 0.001$
Overall Mean	0.696	0.915	+31.5%	$p < 0.001$

Statistical Equivalence Support: The consistent reliability improvements across model variants support the statistical equivalence finding by demonstrating that:

1. **Enhanced Precision:** Improved evaluation reliability reduces measurement error, increasing sensitivity to detect true differences if they exist
2. **Reduced Bias:** Systematic bias mitigation ensures that observed equivalence reflects genuine model performance rather than evaluation artifacts
3. **Consistent Application:** Uniform methodology application across all model variants ensures fair comparison
4. **Methodological Robustness:** High reliability scores indicate that the evaluation methodology provides stable, trustworthy assessment

Cross-Validation of Methodology Effectiveness: Independent validation using hold-out topics confirms methodology effectiveness:

- **Validation Sample:** 15 additional topics not used in primary analysis
- **Consistency Replication:** Hybrid methodology advantages replicated ($p < 0.001$)
- **Effect Size Stability:** Improvement effect sizes remain large (Cohen’s $d > 1.2$)
- **Generalizability:** Benefits extend across different topic categories and complexity levels

Methodology Effectiveness in Supporting Statistical Conclusions

The Hybrid methodology’s enhanced reliability and consistency directly support the validity of the statistical equivalence conclusion by ensuring that the evaluation framework provides accurate, unbiased assessment of model performance.

Statistical Power Enhancement: The methodology improvements translate to enhanced statistical power for detecting meaningful differences:

1. **Reduced Measurement Error:** Higher reliability scores indicate reduced random measurement error, increasing statistical sensitivity
2. **Enhanced Effect Size Detection:** More precise evaluation enables detection of smaller meaningful effects when they exist
3. **Improved Confidence Intervals:** Enhanced reliability produces tighter confidence intervals around effect size estimates
4. **Robust Null Hypothesis Testing:** Systematic bias mitigation ensures that null hypothesis retention reflects true equivalence rather than methodological limitations

Methodology Contribution to Research Validity: The Hybrid methodology represents a significant methodological contribution that enhances the overall validity and reliability of the research conclusions:

- **Internal Validity:** Systematic bias mitigation and consistency verification strengthen causal inference

- **External Validity:** Context-aware evaluation criteria improve generalizability to real-world applications
- **Construct Validity:** Reasoning-enhanced evaluation better captures the intended email quality constructs
- **Statistical Conclusion Validity:** Enhanced reliability and bias mitigation support robust statistical inference

The comprehensive performance validation demonstrates that the Hybrid prompting methodology provides systematic improvements over traditional evaluation approaches, supporting the reliability and validity of the statistical equivalence finding that DPO fine-tuning produces no meaningful improvements in automated email generation tasks.

A.6.5 Technical Implementation and Optimization Procedures

This subsection provides detailed technical specifications for the Hybrid prompting methodology implementation, including coordination algorithms, reasoning model selection criteria, and systematic optimization procedures that ensure reliable, efficient operation across different model configurations and evaluation contexts.

Coordination Algorithms and Agent Management

The Hybrid methodology implements sophisticated coordination algorithms that manage multi-agent interactions while maintaining evaluation integrity and optimizing computational efficiency.

Master Coordination Algorithm: The primary coordination algorithm manages the complete evaluation lifecycle through systematic state management and resource optimization:

Master Coordination Algorithm (MCA):

INPUT: topic_context, model_list, evaluation_config

OUTPUT: aggregated_results, quality_metrics

1. INITIALIZE_SESSION(unique_id, timestamp, config)
 - a. Setup agent configurations and resource allocation
 - b. Initialize state tracking and audit logging
 - c. Validate input parameters and dependencies
2. FOR each model IN model_list:
 - a. GENERATE_EMAIL(model, topic_context)
 - i. Load model with optimized configuration
 - ii. Generate email content with quality validation
 - iii. Store results with metadata tracking
 - b. CREATE_CHECKLIST(topic_context, evaluation_mode)
 - i. Invoke Checklist Creator with reasoning integration

- ii. Validate criteria completeness and relevance
 - iii. Apply priority weighting and category balancing
- c. EVALUATE_WITH_CONSENSUS(email_content, checklist)
 - i. FOR attempt IN [1, 2, 3]:
 - Execute independent Judge Agent evaluation
 - Record binary responses with confidence scores
 - Store detailed evaluation metadata
 - ii. RESOLVE_CONSENSUS(evaluation_attempts)
 - Apply majority voting with confidence weighting
 - Calculate agreement metrics and reliability scores
 - Generate evaluation summary and recommendations
- 3. AGGREGATE_RESULTS(all_model_results)
 - a. Calculate comparative statistics across models
 - b. Generate quality assurance reports
 - c. Validate result consistency and completeness
- 4. FINALIZE_SESSION(results, quality_metrics, audit_trail)

Agent State Management System: The state management system ensures consistent agent behavior and enables comprehensive audit trails:

- **State Persistence:** Comprehensive logging of agent states and transitions
- **Error Recovery:** Automatic retry mechanisms with exponential backoff
- **Resource Management:** Dynamic allocation of computational resources based on demand
- **Performance Monitoring:** Real-time tracking of response times and success rates

Reasoning Model Selection Criteria and Optimization

The Hybrid methodology employs systematic criteria for reasoning model selection and optimization to ensure optimal performance across different evaluation contexts.

Model Selection Framework: The reasoning model selection process follows established criteria optimized for multi-agent coordination and evaluation consistency:

Table A.17: *Reasoning Model Selection and Optimization Criteria*

Selection Criterion	Weight	Measurement	Optimization Target
Reasoning Capability	0.30	Chain-of-thought quality	Maximize logical coherence
Consistency Score	0.25	Inter-evaluation agreement	Minimize variance
Context Integration	0.20	Topic relevance accuracy	Maximize contextual alignment
Computational Efficiency	0.15	Response time and resources	Minimize resource usage
Bias Mitigation	0.10	Fairness across categories	Maximize evaluation equity

Optimization Procedures: The reasoning model optimization follows systematic procedures to enhance performance while maintaining evaluation reliability:

- 1. **Performance Profiling:** Systematic evaluation of model performance across validation topics:
 - Reasoning quality assessment using expert human evaluation
 - Consistency measurement through repeated evaluation protocols
 - Context integration verification through topic-specific performance analysis
- 2. **Parameter Optimization:** Fine-tuning of model parameters for optimal evaluation performance:
 - Temperature optimization: Systematic testing across range [0.1, 0.5]
 - Token limits: Optimal balance between reasoning depth and efficiency
 - Sampling parameters: Reproducibility vs. diversity optimization
- 3. **Configuration Validation:** Comprehensive validation of optimized configurations:
 - Cross-validation using holdout topics
 - Stress testing under high-load conditions
 - Stability assessment across different computational environments

Performance Monitoring and Quality Assurance Procedures

The technical implementation includes comprehensive performance monitoring and quality assurance procedures that ensure consistent operation and enable continuous improvement.

Real-Time Performance Monitoring: The system implements continuous performance monitoring across multiple dimensions:

Table A.18: Real-Time Performance Monitoring Metrics

Performance Metric	Target	Warning Threshold	Critical Threshold	Action
Agent Response Time	< 60s	> 45s	> 75s	Auto-retry/Scale
Evaluation Agreement	> 0.85	< 0.80	< 0.70	Expert Review
Memory Utilization	< 80%	> 75%	> 90%	Resource Scaling
Error Rate	< 2%	> 3%	> 5%	System Alert
Queue Processing Rate	> 0.8 req/min	< 0.6 req/min	< 0.4 req/min	Load Balancing

Automated Quality Assurance Pipeline: The quality assurance system operates continuously to ensure methodology integrity:

Quality Assurance Pipeline (QAP):

- 1. DATA_VALIDATION:
 - Schema compliance checking for all inputs/outputs
 - Completeness verification across evaluation components

- Integrity validation using cryptographic checksums
2. CONSISTENCY_MONITORING:
 - Real-time agreement ratio calculation
 - Statistical outlier detection and flagging
 - Cross-topic performance variance monitoring
 3. BIAS_DETECTION:
 - Systematic bias pattern analysis using statistical tests
 - Fairness metrics calculation across topic categories
 - Evaluation drift detection through control questions
 4. PERFORMANCE_OPTIMIZATION:
 - Response time monitoring and optimization
 - Resource utilization analysis and scaling
 - Error rate tracking and preventive maintenance
 5. REPORTING_AND_ALERTS:
 - Automated generation of quality reports
 - Real-time alerting for threshold violations
 - Trend analysis and predictive maintenance recommendations

Scalability and Deployment Considerations

The Hybrid methodology implementation includes comprehensive scalability provisions and deployment optimization for research and production environments.

Horizontal Scaling Architecture: The system supports horizontal scaling through distributed agent deployment:

- **Agent Distribution:** Independent agent instances enabling parallel processing
- **Load Balancing:** Intelligent request distribution across available resources
- **State Synchronization:** Consistent state management across distributed components
- **Fault Tolerance:** Automatic failover and recovery mechanisms

Resource Optimization Strategies: Systematic optimization ensures efficient resource utilization:

1. **Memory Management:** Dynamic memory allocation with automatic garbage collection
2. **Computational Optimization:** GPU utilization optimization for model inference
3. **Network Efficiency:** Minimized data transfer through intelligent caching
4. **Storage Optimization:** Compressed data formats and efficient indexing

Deployment Configuration Management: The implementation includes comprehensive configuration management for different deployment scenarios:

Configuration Management Framework:

- DEVELOPMENT: Single-node deployment with full logging and debugging
- TESTING: Multi-node deployment with comprehensive validation
- STAGING: Production-like environment with performance monitoring
- PRODUCTION: Optimized deployment with minimal logging overhead
- RESEARCH: Enhanced auditing and detailed performance analytics

Technical Implementation Summary: The Hybrid prompting methodology technical implementation provides a robust, scalable, and optimized framework for multi-agent evaluation that enables reliable research conclusions while maintaining high performance and quality standards. The comprehensive technical specifications ensure reproducibility, maintainability, and extensibility for future research applications and practical deployments.